Concurrent Data Processing in Elixir

Fast, Resilient Applications with OTP, GenStage, Flow, and Broadway

Svilen Gospodinov

Foreword: *José Valim*
Series editor: *Bruce A. Tate*
Development editor: *Jacquelyn Carter*

Under Construction: The book you're reading is still under development. As part of our Beta book program, we're releasing this copy well before a normal book would be released. That way you're able to get this content a couple of months before it's available in finished form, and we'll get feedback to make the book even better. The idea is that everyone wins!

Be warned: The book has not had a full technical edit, so it will contain errors. It has not been copyedited, so it will be full of typos, spelling mistakes, and the occasional creative piece of grammar. And there's been no effort spent doing layout, so you'll find bad page breaks, over-long code lines, incorrect hyphenation, and all the other ugly things that you wouldn't expect to see in a finished book. It also doesn't have an index. We can't be held liable if you use this book to try to create a spiffy application and you somehow end up with a strangely shaped farm implement instead. Despite all this, we think you'll enjoy it!

Download Updates: Throughout this process you'll be able to get updated ebooks from your account at *pragprog.com/my_account*. When the book is complete, you'll get the final version (and subsequent updates) from the same address.

Send us your feedback: In the meantime, we'd appreciate you sending us your feedback on this book at *pragprog.com/titles/sgdpelixir/errata*, or by using the links at the bottom of each page.

Thank you for being part of the Pragmatic community!

The Pragmatic Bookshelf

Concurrent Data Processing in Elixir

Fast, Resilient Applications with OTP, GenState, Flow, and Broadway

Svilen Gospodinov

Pragmatic Bookshelf

Many of the designations used by manufacturers and sellers to distinguish their products are claimed as trademarks. Where those designations appear in this book, and The Pragmatic Programmers, LLC was aware of a trademark claim, the designations have been printed in initial capital letters or in all capitals. The Pragmatic Starter Kit, The Pragmatic Programmer, Pragmatic Programming, Pragmatic Bookshelf, PragProg and the linking *g* device are trademarks of The Pragmatic Programmers, LLC.

Every precaution was taken in the preparation of this book. However, the publisher assumes no responsibility for errors or omissions, or for damages that may result from the use of information (including program listings) contained herein.

For our complete catalog of hands-on, practical, and Pragmatic content for software developers, please visit *https://pragprog.com*.

For sales, volume licensing, and support, please contact *support@pragprog.com*.

For international rights, please contact *rights@pragprog.com*.

Contents

Introduction

Data processing is an essential part of many software applications. In fact, most engineers don't even think about it as something separate to programming. But if you're transforming information in some way, for example, when doing reporting, data aggregation, or analytics, then you're doing data processing.

Thanks to the Erlang Virtual Machine (also known as the BEAM), everyone who uses Elixir benefits from it's amazing concurrency model, which is particularly well suited for long-running, concurrent tasks. As a result, you will find that Elixir offers more ways for performing concurrent work than other languages.

While this is a good thing, it could be also challenging to find the right tool for the job. Some tasks are excellent fit for Flow, while others are perfect for Broadway. Sometimes it's easier to just use GenStage, or even the Task module. Making the right choice means going through dozens of pages of documentation, and that would be just the beginning.

This book is here to help you navigate the world of concurrency tools available in the Elixir ecosystem. You will learn about the most popular modules and libraries, and start using them in no time. You will also discover a range new of new techniques, that will help you simplify your product, improve the performance of your code, and make your application more resilient to errors and increased workloads.

Who Should Read This Book?

Most applications have to process data in some way, and run on machines with multi-core CPUs. If your application is one of them, and data processing plays an important part in what you do, then this book is for you. Web development is very popular with Elixir, but the techniques in this book work equally well if you're using Nerves and working on embedded software.

This book is targeted at intermediate developers, who are comfortable with writing Elixir code. If you're not familiar with the language, then you should pick another book first, before attempting this one.

However, you don't have to be an expert in Elixir—this book will guide you through each topic with plenty of examples and helpful figures along the way. After every chapter, you will gain valueable knowledge which you can then apply to your personal or business projects.

About This Book

There are five chapters, each dedicated to a specific module or a library.

In Chapter 1, Easy Concurrency With the Task Module, on page 1, you'll get started on the journey of learning how concurrency works in Elixir. It introduces the Task module, processes, timeouts, and other topics that lay the foundation for the following chapters to built upon.

In Chapter 2, Long-Running Processes Using GenServer, on page 25 you'll learn about GenServer and supervisors. You will see how to create and configure GenServer processes by building a simple job processing system. We'll introduce the Supervisor behaviour and talk about how Elixir achieves fault-tolerance.

In Chapter 3, Data Processing Pipelines with GenStage, on page 57 we'll move on to data processing pipelines. You will learn about back-pressure, and the building blocks of GenStage—producer, consumer, and producer-consumer. You will also start building your very own web scraper by putting what you have learned in practice.

In chapter 4, we'll cover Flow, and how you can use it instead of GenStage for operations like map, filter, reduce, and more. You will see how you can use Flow when working with large datasets, and even plug it into existing GenStage data processing pipelines.

In the last chapter, we'll explore Broadway. We're going to set up a data-ingestion pipeline using RabbitMQ, but the techniques apply to other message brokers, such as Amazon SQS, Apache Kafka, Google Cloud Pub/Sub. We will talk cover the various options and benefits that come with Broadway.

About the Code

You can't apply all the techniques in this book, without having data to process or services to integrate with. At the same time, downloading large data sets, or signing up to third-party services, is too cumbersome and not practical for the purposes of this book. That's why all projects attempt to simulate real-

world cases, so you can focus on the implementation details. It also makes them easy to reproduce.

To run the code examples and snippets from the book on your computer, make sure you have Erlang and Elixir installed first. The versions used in this book are Erlang 23.0 and Elixir 1.11. If you need help installing them, check Elixir's official website[1] to see what installation options you have for your operating system.

For Chapter 5, you will also need RabbitMQ installed. You can find installation instructions on RabbitMQ's website.[2]

Online Resources

All examples and source code can be downloaded from the book page on the Pragmatic Bookshelf website.[3] Please report any errors or suggestions you may have, using the errata link.[4]

If you enjoyed reading this book, please let others know about it and spread the word. If you use Twitter, you can find me at @svileng[5] and let me know which part of the book you liked the most. You can also tweet to @pragprog.[6]

Svilen Gospodinov

February 2021

---

1. https://elixir-lang.org/install.html
2. https://www.rabbitmq.com/download.html
3. https://pragprog.com/book/sgdpelixir
4. https://pragprog.com/book/sgdpelixir/errata
5. https://twitter.com/svileng
6. https://twitter.com/pragprog

Easy Concurrency With the Task Module

Since the dawn of the computer industry, hardware manufacturers and computer scientists have tried to make computers faster at running programs. At first, multi-threading was the only way to achieve concurrency, which is the ability to run two or more programming tasks, and switch between them to collect the results. This is how computers appeared to be doing many things at once, when in fact they were simply multitasking.

Multi-core CPUs changed that. They brought parallelism and allowed tasks to run at the same time, independently, which significantly increased systems' performance. Multi-processor architectures followed, enabling even greater concurrency and parallelism by supporting two or more CPUs on a single machine. The figure below shows a simple comparison between concurrency on a single-core CPU and a dual-core CPU. The latter also enables parallelism:
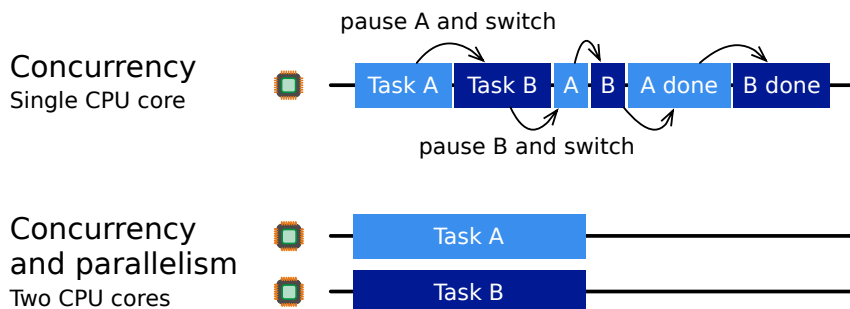
Of course, cutting-edge hardware always comes with a high price tag. But with the advent of cloud computing, things changed once again. Nowadays

you can run code on cloud services using virtual machines with dozens of CPU cores, without the need to buy and maintain any physical hardware.

All these advancements are important to us, software engineers. We want to write software that performs well and runs quickly. After all, no one likes loading screens and waiting for the computer to finish. However, running code on a multi-core processor system does not automatically make it efficient. In order to take full advantage of the computer resources available to us, we need to write software with concurrency and parallelism in mind. Thankfully, modern programming languages try to help us as much as possible, and Elixir is no exception. In fact, thanks to Erlang, the Erlang Virtual Machine (BEAM), and the Open Telecom Platform (OTP), Elixir is a superb choice for building concurrent applications and processing data, as you'll see soon in this and upcoming chapters.

In this book we're going to cover the most popular tools for performing concurrent work using Elixir. You will learn about the pros and cons of each one, and see how they work in practice. Some of them, like the Task module and GenServer, come with Elixir. The others—GenStage, Flow, and Broadway—are available as stand-alone libraries on the Hex.pm package registry. Knowing how to utilize each of these tools will help you leverage concurrency in the most effective way, and solve even the most challenging problems. Along the way, you will also learn how to build fault-tolerant applications, recover from failures, use back-pressure to deal with limited system resources, and many more useful techniques.

First, we are going to look at the Task module, which is part of the Elixir standard library. It has a powerful set of features that will help you run code concurrently. You are also going to see how to handle errors and prevent the application from crashing when a concurrent task crashes. The chapter provides a foundation on which the following chapters will be built upon, so let's get started!

Introducing the Task Module

To run code concurrently in Elixir, you have to start a process and execute your code within that process. You may also need to retrieve the result and use it for something else. Elixir provides a low-level function and a macro for doing this—spawn/1 and receive. However, using them could be tricky in practice, and you will likely end up with a lot of repetitive code.

Elixir also ships with a module called Task, which significantly simplifies starting concurrent processes. It provides an abstraction for running code

concurrently, retrieving results, handling errors and starting a series of processes. It packs a lot of features and has a concise API, so there is rarely (if ever) need to use the more primitive spawn/1 and receive.

In this chapter, we are going to cover everything that the Task module has to offer. You will learn how to start tasks, and different ways to retrieve results. You will tackle processing large lists of data. We will talk about handling failure, and explain how process linking works in Elixir. You will then see how to use one of the built-in Supervisor modules for isolating process crashes, and finally, discuss Elixir's approach to error handling.

Before we dive in, let's create an Elixir project to work on first, and get familiar with some of the development tools we're going to use throughout this and the following chapters.

What is an Elixir process?

Processes in Elixir are Erlang processes, since Elixir runs on the Erlang Virtual Machine. Unlike operating system processes, they are very lightweight in terms of memory usage and quick to start. The Erlang VM knows how to run them concurrently and in parallel (when a multi-core CPU is present). As a result, by using processes, you get concurrency and parallelism for free.

Creating Our Playground

We are going to create an application called sender and pretend that we are sending emails to real email addresses. We are going to use the Task module later to develop some if its functionality.

First, let's use the mix command-line tool to scaffold our new project:

```
$ mix new sender --sup
```

This creates a sender directory with a bunch of files and folders inside. Notice that we also used the --sup argument, which will create an application with a *supervision tree*. You will learn about supervision trees later in this chapter.

Next, change your current directory to sender with cd sender and run iex -S mix. You should see some Erlang version information and the following message:

```
Interactive Elixir (1.9.4) - press Ctrl+C to exit (type h() ENTER for help)
iex(1)>
```

We're now running the *Interactive Elixir* shell, also known as IEx. We are going to use it to test code frequently throughout the book. Most of the time, when

we make a code change using our text editor, we can call the special recompile/0 function available in IEx, and the code will reload:

```
iex(1)> recompile()
:noop
```

We haven't actually added or changed any code yet, so the function returned just :noop for *no operation*.

In rare cases, you may need to restart the IEx shell entirely, for example when making fundamental application changes in the application.ex file. You can restart IEx by pressing Ctrl-C twice to quit, and then running the iex -S mix command again.

To keep our project and examples simple, we're not actually going to send real emails. However, we still need some business logic for our experiments. We can use the Process.sleep/1 function to pretend we're sending an email, which is normally a slow operation and can take a few seconds to complete. When called with an integer, Process.sleep/1 stops the current process for the given amount of time in milliseconds. This is very handy, because you can use it to simulate code that takes a long time to complete. You can also use it to test various edge cases, as you will see later. Of course, in real world production applications, you will replace this with your actual business logic. But for now, let's pretend that we're doing some very intensive work.

Let's open sender.ex and add the following:

```
sender/lib/sender.ex
def send_email(email) do
  Process.sleep(3000)
  IO.puts("Email to #{email} sent")
  {:ok, "email_sent"}
end
```

Calling this function will pause execution for three seconds and print a message, which will be useful to debugging. It also returns a tuple {:ok, "email_sent"} to indicate that the email was successfully sent.

Now that everything is set up we're ready to start. I suggest you keep one terminal session with IEx open and your favorite text editor next to it, so you can make and run changes as we go.

Starting Tasks and Retrieving Results

Before jumping into the Task module, let's see how things work at the moment. Let's call the send_email/1 function from IEx, and pass a fictional email address

as an argument. Don't forget to run the recompile() command first, if you've been running IEx already:

```
iex> recompile()
Compiling 1 file (.ex)
:ok

iex> Sender.send_email("hello@world.com")
Email to hello@world.com sent
{:ok, "email_sent"}
```

Did you notice the delay? We had to wait three seconds until we saw the printed output and result. In fact, even the iex> prompt was not showing. Let's add another function, notify_all/1:

sender/lib/sender.ex
```
def notify_all(emails) do
  Enum.each(emails, &send_email/1)
end
```

The notify_all/1 function uses Enum.each/2 to iterate over the variable emails, which is a list of strings. For each item in the list, we are going to call the send_email/1 function. We are going to test this function in IEx, but first we need some test data. Create a file .iex.exs in the main project folder sender, at the top level where mix.exs is also located. Add the following:

sender/.iex.exs
```
emails = [
  "hello@world.com",
  "hola@world.com",
  "nihao@world.com",
  "konnichiwa@world.com",
]
```

Once you save the file, quit IEx and start it again with iex -S mix. Type emails and press enter to inspect the variable:

```
iex(1)> emails
["hello@world.com", "hola@world.com", "nihao@world.com",
 "konnichiwa@world.com"]
```

This will save you a lot of typing. All Elixir code in .iex.exs will run when IEx starts so this will persist between IEx sessions. Now let's use the test data with notify_all/1. Can you guess how much time it will take to send all emails? Let's find out:

```
iex> Sender.notify_all(emails)
Email to hello@world.com sent
Email to hola@world.com sent
Email to nihao@world.com sent
```

```
Email to konnichiwa@world.com sent
:ok
```

It took four calls to send_email/1 and whole fourteen seconds to complete. Just waiting for the output in IEx felt like ages. This is not good at all. As our user base grows, it will take forever to send our emails.

Don't despair—we can significantly speed up our code using the Task module. However, before we jump into it, let's take a moment to talk about two important concepts in programming: *synchronous* and *asynchronous* code.

Synchronous and Asynchronous Code

By default, when you run some code in Elixir, you have to wait for it to complete. You cannot do anything in the meantime and you get the result as soon as code has finished. The code is executed *synchronously* and is sometimes called *blocking code*.

The opposite of this is running code *asynchronously*. In this case, you ask the programming runtime to run the code, but carry on with the rest of the program. Asynchronous code runs in the *background* because the application keeps running, as if nothing happened. Eventually, when the asynchronous code finishes, you can retrieve the result. Because asynchronous code does not block the main execution of the program, it is called *non-blocking code*. Asynchronous code is also concurrent, since it enables us to continue doing other work. The figure below illustrates the difference between synchronous and asynchronous code in a bread-making program:
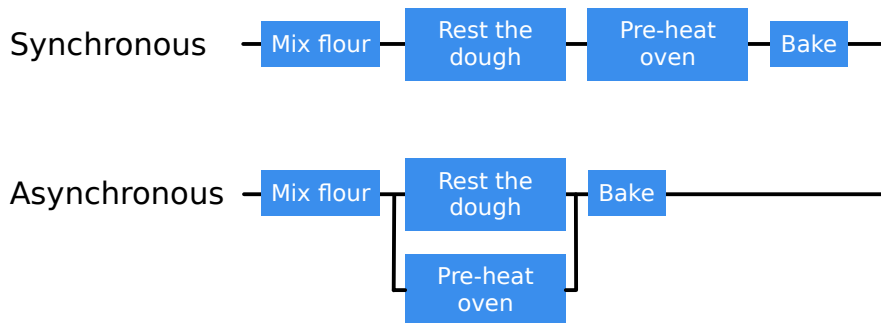
Since only the final step Bake requires the oven to be pre-heated, you can run Pre-heat oven asynchronously and do something else in the meantime. Compared

to the synchronous version, this significantly decreases the time needed to complete the set of instructions.

You may have guessed already that the notify_all/1 function is sending each email synchronously. This is why it is taking so long to complete. To improve this, we're going to convert our code to run asynchronously. Thanks to Elixir's Task module we will need only a few small changes to achieve that. Let's see how it works.

Starting Processes

The Task module contains a number of very useful functions for running code asynchronously and concurrently. One of them is start/1. It accepts a function as an argument, and inside that function we should do all the work we intend to do. Let's try it out quickly in IEx:

```
iex> Task.start(fn -> IO.puts("Hello async world!") end)
Hello async world!
{:ok, #PID<0.266.0>}
```

You saw the message printed instantly because it's a quick operation, but the actual result was {:ok, #PID<0.266.0>}. Returning a tuple like {:ok, result} or {:error, message} is a common practice in Elixir. The result was :ok for success and #PID<0.266.0>. *PID* stands for *process identifier*—a number that uniquely identifies an Elixir process.

We already have send_email/1 ready, so we can use Task.start/1 to call it. Let's make some changes:

```
sender/lib/sender.change1.ex
def notify_all(emails) do
  Enum.each(emails, fn email ->
    Task.start(fn ->
      send_email(email)
    end)
  end)
end
```

And then recompile and run notify_all/1:

```
iex> Sender.notify_all(emails)
:ok
Email to hello@world.com sent
Email to hola@world.com sent
Email to nihao@world.com sent
Email to konnichiwa@world.com sent
```

This should be significantly faster—in fact, four times faster! All functions were called concurrently and finished at the same time, printing the success message as we expected.

Retrieving the Result of a Task

Task.start/1 has one limitation by design: it does not return the result of the function that was executed. This may be useful in some cases, but most of the times you need the result for something else. It would be great if we modify our code and return a meaningful result when all emails are sent successfully.

To retrieve the result of a function, you have to use Task.async/1. It returns a %Task{} struct which you can assign to a variable for later use. You can give it a try in IEx like so:

```
iex> task = Task.async(fn -> Sender.send_email("hello@world.com") end)
%Task{
  owner: #PID<0.145.0>,
  pid: #PID<0.165.0>,
  ref: #Reference<0.713486762.1657274369.63141>
}
```

The send_email/1 code is now running in the background. In the meantime, we are free to do other work as we need. You can add more business logic or even start other tasks. When you need the actual result of the task, you can retrieve it using the task variable. Let's have a closer look at what this variable contains:

- owner is the PID of the process that started the Task process;
- pid is the identifier of the Task process itself;
- ref is the *process monitor reference*.

Process monitoring is out of the scope for this book. However, it is worth knowing that you can monitor a process and receive notifications from it using a reference value—for example, when and how a process exits.

To retrieve the result of the task, you can use either Task.await/1 or Task.yield/1 which accept a Task struct as an argument. There is an important difference in the way await/1 and yield/1 work, so you have to choose wisely. They both stop the program and try to retrieve the result of the task. The difference comes from the way they handle *process timeouts*.

Process timeouts ensure that processes don't get stuck running forever. To show how they work, let's increase the time the send_email/1 function takes to 30 seconds, just temporarily:

```
def send_email(email) do
  Process.sleep(30_000)
  IO.puts("Email to #{email} sent")
  {:ok, "email_sent"}
end
```

And recompile/0 in IEx. We're now going to run the code asynchronously and pipe the task result into await/1:

```
iex> Task.async(fn -> Sender.send_email("hi@world.com") end) |> Task.await()
```

After five seconds, you will receive an exception similar to this one:

```
** (exit) exited in: Task.await(%Task{owner: #PID<0.144.0>,
  pid: #PID<0.151.0>,
  ref: #Reference<0.2297312895.3696492546.156249>}, 5000)
    ** (EXIT) time out
    (elixir) lib/task.ex:607: Task.await/2
```

When using await/1 we expect a task to finish within a certain amount of time. By default, this time is set to 5000ms, which is five seconds. You can change that by passing an integer with the amount of milliseconds as a second argument, for example Task.await(task, 10_000). You can also disable the timeout by passing the atom :infinity.

In comparison, Task.yield/1 simply returns nil if the task hasn't completed. The timeout of yield/1 is also 5000ms but does not cause an exception and crash. You can also do Task.yield(task) repeatedly to check for a result, which is not allowed by await/1. A completed task will return either {:ok, result} or {:exit, reason}. You can see this in action:

```
iex> task = Task.async(fn -> Sender.send_email("hi@world.com") end)
%Task{
  owner: #PID<0.135.0>,
  pid: #PID<0.147.0>,
  ref: #Reference<0.3033103973.1551368196.24818>
}
iex> Task.yield(task)
nil
iex> Task.yield(task)
nil

Email to hi@world.com sent
iex> Task.yield(task)
{:ok, {:ok, "email_sent"}}
```

This output illustrates what happens when you call Task.yield/1 while the task is still running—you receive nil. Once the success message is printed out, we

received {:ok, {:ok, "email_sent"}} as expected. You can also use yield/2 and provide your own timeout, similarly to await/2, but the :infinity option is not allowed.

You may be wondering what happens if our task is stuck and never finishes? While await/1 takes care of stopping the task, yield/1 will leave it running. It is a good idea to stop the task manually by calling Task.shutdown(task). The shutdown/1 function also accepts a timeout and gives the process a last chance to complete, before stopping it. If it completes, you will receive the result as normal. You can also stop a process immediately (and rather violently) by using the atom :brutal_kill as a second argument.

As you can see, using yield/1 and shutdown/1 is a bit more work than await/1. Which one to use depends on your use case. Very often a task timeout justifies an exception, in which case await/1 will be more convenient to use. Whenever you require more control over the timeout and shutdown, you can switch to yield/1 and shutdown/1.

For our notify_all/1 logic, we're going to use await/1 for simplicity. Remember to revert our previous change in send_email/1 and set Process.sleep/1 back to 3000ms:

```
def send_email(email) do
  Process.sleep(3000)
```

Now we're going to replace Task.start/1 with Task.async/1:

```
sender/lib/sender.change2.ex
def notify_all(emails) do
  emails
  |> Enum.map(fn email ->
    Task.async(fn ->
      send_email(email)
    end)
  end)
  |> Enum.map(&Task.await/1)
end
```

Notice that we're also using Enum.map/2 instead of Enum.each/2, because we want to map each email string to its corresponding Task struct. The struct is needed to retrieve the result from each process.

We used the Elixir shorthand function syntax here—&Task.await/1. If you are not familiar with it, it is simply equivalent to writing:

```
Enum.map(fn task ->
  Task.await(task)
end)
```

Let's try out the latest changes in IEx:

```
iex> Sender.notify_all(emails)
Email to hello@world.com sent
Email to hola@world.com sent
Email to nihao@world.com sent
Email to konnichiwa@world.com sent
[ok: "email_sent", ok: "email_sent", ok: "email_sent", ok: "email_sent"]
```

The function returned a list of results. For each task you have an ok: "email_sent" tuple, which is what the send_email/1 function returns. In your business logic, you can return an {:error, "error message"} when something goes wrong. Then you will be able to collect the error and potentially retry the operation or do something else with the result.

Creating tasks from lists of items is actually very common in Elixir. In the next section, we are going to use a function specifically designed for doing this. It also offers a range of additional features, especially useful when working with large lists. Keep reading.

Keyword Lists in Elixir

Elixir has a special notation for lists containing key-value tuples, also known as keyword lists.[1] Each item in the keyword list must be a two-element tuple. The first element is the key name, which must be an atom. The second is the value and can be of any type. They are displayed in IEx without the curly braces, like this [ok: "email_sent"]. This is equivalent to [{:ok, "email_sent"}].

Managing Series of Tasks

Let's imagine we have one million users and we want to send an email to all of them. We can use Enum.map/2 and Task.async/1 like we did before, but starting one million processes will put sudden pressure on our system resources. It can degrade system's performance and potentially make other services unresponsive. Our email service provider will not be happy either, because we also put a lot of pressure on their email infrastructure.

On the other hand, we don't want to send emails one by one, because it is slow and inefficient. It seems that we are at a crossroad, and whichever way we take, we end up in peril. We don't want to choose between performance and reliability—we want to be able to run Task processes to leverage concurrency, but ensure we do not overload our system resources, as we scale our product and increase our user base.

---

1.   https://elixir-lang.org/getting-started/keywords-and-maps.html

The solution to our problems is async_stream/3. It's another very handy function from the Task module which is designed to create task processes from a list of items. For every item in the list, async_stream/3 will start a process and run the function we provide to process the item. It works just like Enum.map/2 and Task.async/2 combined, with one major difference: you can set a limit on the number of processes running at the same time. The figure below illustrates how this works:
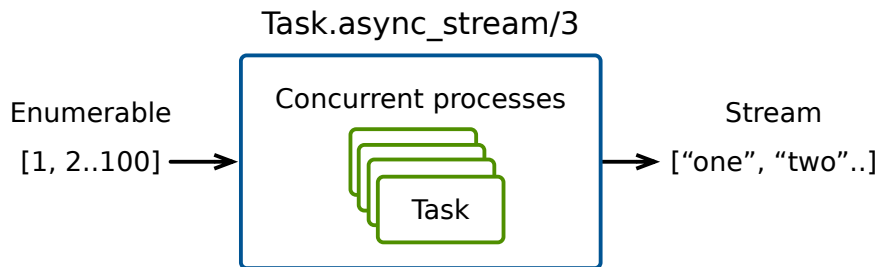
In this example, the concurrency limit is set to four, so even if you have a list of one hundred items, only four processes at most will run concurrently at any given time. This is an example of handling *back-pressure*, which we are going to discuss in-depth in . For now, all you need to know is that this strategy of handling processes is great at preventing sudden spikes of system usage. In other words, you get concurrency and performance, without sacrificing reliability.

As the function name suggests, async_stream/3 returns a Stream. Streams in Elixir are data structures that hold one or more operations that don't run immediately, only when explicitly told so. That's why they're sometimes called *lazy enumerables*. Let's see what happens when we run this function from IEx:

```
iex> Task.async_stream(emails, &Sender.send_email/1)
#Function<1.35903181/2 in Task.build_stream/3>
```

Instead of the usual result, we received a function, which is going to create a Stream. It's important to understand how streams work, so let's run another quick example in IEx:

```
iex> Stream.map([1, 2, 3], & &1 * 2)
#Stream<[
  enum: [1, 2, 3],
  funs: [#Function<49.33009823/1 in Stream.map/2>]
```

```
]>
```

If we have used Enum.map/2 the function would have returned [2, 4, 6]. Instead, the result stream simply contains the initial input and a list of operations funs. These operations can be run at a later time. Because both Stream and Enum implement the Enumerable protocol, many Enum functions have a lazy alternative in the Stream module.

One way to run a stream is to use the Stream.run/1 function. However, Stream.run/1 always returns :ok so it is only useful when you are not interested in the final result. Instead, you can use Enum.to_list/1 which will try to convert the stream to a List data structure. As a result of this conversion, all operations in the stream will run and the result will be returned as a list. Other functions in the Enum module will also force the stream to run, such as Enum.reduce/3. You can use them if you intend to do more work with the result.

Why does async_stream return a stream?

Streams are designed to emit a series of values, one by one. As soon as Task.async_stream/3 finds out that a Task process has completed, it will emit the result and take the next element from the input, starting a new process. This means it can maintain a number of concurrent events, which is one of the benefits of async_stream/3. You can also use all other functions from the Stream module to compose complex data processing flows.

Now, let's update notify_all/1 to use async_stream/3. This time, however, we will run the stream using Enum.to_list/1:

sender/lib/sender.change3.ex
```
def notify_all(emails) do
  emails
  |> Task.async_stream(&send_email/1)
  |> Enum.to_list()
end
```

And give it a go in IEx, but don't forget to recompile/0 first:

```
iex> Sender.notify_all(emails)
Email to hello@world.com sent
Email to hola@world.com sent
Email to nihao@world.com sent
Email to konnichiwa@world.com sent
[
  ok: {:ok, "email_sent"},
  ok: {:ok, "email_sent"},
  ok: {:ok, "email_sent"},
  ok: {:ok, "email_sent"}
```

```
]
```

As you can see, the output is similar to the one from the Task.async/2 example. However, depending on how many logical cores your machine has, the time it takes for the function to complete may be different.

As we mentioned before, async_stream/3 maintains a limit on how many processes can be running at the same time. By default, this limit is set to the number of logical cores available in the system. Previously we used Task.async/2 to manually start a process for each of the four items. This means that if you have a CPU with less than four logical cores, async_stream/3 will appear to be slower. You can easily change this default behavior through the optional max_concurrency parameter. Let's set max_concurrency to 1 temporarily:

```
|> Task.async_stream(&send_email/1, max_concurrency: 1)
```

When you try the new changes again, you will see that emails are sent out one by one. This is not very useful in practice, but it demonstrates how max_concurrency works. You can revert the change or set it to an even higher number. In production applications, you can benchmark if higher max_concurrency works better for your use case, taking into consideration your system resources and need for performance.

Another option that needs mentioning is :ordered. Currently, async_stream/3 assumes that we want the results in the same order as they were originally. This order preservation can potentially slow down our processing, because async_stream/3 will wait for a slow process to complete, before moving on to the next.

In our case, we only need the results to check if an email was successfully sent or not. We don't necessarily need them in exactly the same order. We can potentially speed things up by disabling ordering like so:

```
|> Task.async_stream(&send_email/1, ordered: false)
```

Now async_stream/3 won't be idle if one process is taking longer than others.

Processes started by async_stream/3 are also subject to timeouts, just like those started by start/1 and async/2. The :timeout optional parameter is supported and defaults to 5000ms. When a task reaches the timeout, it will produce an exception, stopping the stream and crashing the current process. This behavior can be changed using an optional :on_timeout argument, which you can set to :kill_task. This argument is similar to the :brutal_kill one supported by Task.shutdown/2. Here is an example:

```
|> Task.async_stream(&send_email/1, on_timeout: :kill_task)
```

As a result of using the :kill_task option, when a process exits with a timeout, async_stream/3 will ignore it and carry on as normal.

A process can crash for many reasons. We talked about timeout exceptions, which happen when the processes takes too long to complete. This could be because the work it is doing is very time consuming or because it is waiting for a slow third party API to respond. Sometimes it is an unexpected error that causes an exception. The important thing is that when a process crashes, it can also crash the process that started it, which in turn crashes it's parent process, triggering a chain reaction that can ultimately crash your whole application.

This sounds like a programming disaster waiting to happen, but the good news is that Elixir, thanks to Erlang, has a set of powerful tools to manage processes, catch crashes and recover quickly. In the next section, we are going to explain why this chain reaction is a good thing, how to isolate it and use it to our advantage. You will also start learning how to build the fault-tolerant applications Elixir is famous for.

Linking Processes

Processes in Elixir can be linked together, and Task processes are usually linked to the process that started them automatically. This link is simply called a *process link*. Process links have an important role when building concurrent and fault-tolerant applications. They help us immediately shut down parts of the system or even the whole system when needed, preventing the application from running with a bad state. Let's see why process links are useful and how they work, by looking at a few examples. The following figure shows three processes linked together:


As you can see, Process C has just crashed. Imagine that Process B continues running, unaware of the crash. It is expecting to receive a result from Process C, and report back to Process A. However, this will never happen, so this system is left in a bad state.

This can be avoided thanks to process linking. When two processes are linked, they form a special relationship—as soon as one exits, the other will be notified. When you have a chain of linked processes, all of them will be eventually notified when a crash happens. By default, linked processes will terminate and clean up the memory they use, preventing other potential problems down the line. This means that the crash of Process C will trigger a chain reaction and Process B will terminate, followed by Process A.

However, a chain reaction like this only makes sense when a severe error happens. Take a look at this figure, which shows running processes for a banking web application:

Most web applications rely heavily on a database, and access to the database is usually managed by a process. In this example, the Repo process has this responsibility. If Repo crashes, because the database becomes unavailable, this application won't be able to work at all. If that ever happens, you can let all processes terminate, which will bring the website offline.

Now, let's consider the process Sender, which is responsible for sending emails. Unlike the Repo process, it is not vital in giving users access to their bank account. If the Sender process crashes for whatever reason, we don't want to shut down the whole application. We want to contain this crash and let the rest of the system carry on.

You can isolate crashes by configuring a process to *trap exits*. Trapping an exit means acknowledging the exit message of a linked process, but continue

running instead of terminating. This also means that the exit message won't be propagated further to other processes. Going back to our original example—if Process B is configured to trap exits, it will continue working after the crash of Process C, keeping Process A safe as well. You can configure a process to trap exits manually, but usually you want to use a special type of process called a *supervisor*. We are going to talk about supervisors in just a moment.

As you can see, process links are an essential mechanism for detecting exits and dealing with unexpected process crashes. Most utility functions in the Task module create process links automatically, linking to the current process, but there are some that don't. You have the choice to link or not to link to the current process, but you have to choose the right function.

For example, when we used async/1 and async_stream/3, a process link was created for each new process. Task.start/1, on the other hand, does not create a process link, but there is Task.start_link/1 that does just that.

Most functions in the Task module that link to the current process by default have an alternative, usually ending with _nolink. Task.Supervisor.async_nolink/3 is the alternative to Task.async/1. Task.async_stream/3 can be replaced with Task.Supervisor.async_stream_nolink/4. All functions in the Task.Supervisor module are designed to be linked to a supervisor.

Next, we're going to learn about supervisors and how we can use one when starting task processes.

Meeting the Supervisor

Just like in the workplace, where supervisors are held responsible for groups of employees, Elixir supervisors are responsible for the processes assigned to them. Subordinate processes are also required to report back to their supervisor, who has to ensure everything is running smoothly. To accomplish this, supervisors come with a set of features that allow them to effectively manage other processes. They can start and stop processes, and restart them in case of unforeseen errors in the system. They are configured to trap exits, so when a supervised process exits with an error, that error will be isolated and it won't propagate further. All this makes supervisors an important building block for fault-tolerant applications.

Supervised processes are called *child processes*. Any OTP process can be supervised and you can also add a supervisor as a child of another supervisor. All you have to do is ask the supervisor to start the process you want to be managed. This allows us to easily build a hierarchy of processes also called a *supervision tree*.

Our sender application already has a supervisor in place, which you can see in application.ex:

```
sender/lib/sender/application.ex
def start(_type, _args) do
  children = [
    # Starts a worker by calling: Sender.Worker.start_link(arg)
    # {Sender.Worker, arg}
  ]

  # See https://hexdocs.pm/elixir/Supervisor.html
  # for other strategies and supported options
  opts = [strategy: :one_for_one, name: Sender.Supervisor]
  Supervisor.start_link(children, opts)
end
```

The start/2 function has some boilerplate code and some instructions already in place. At the end of the function we have Supervisor.start_link(children, opts) that starts the main supervisor of the application. Since the children variable is just an empty list, there are actually no child processes to supervise. This setup is the result of us using the --sup argument when calling the mix new command to create our project.

Now that you know how useful supervisors are, let's see how we can use one for our Task processes. We don't want our current process to crash if a single task crashes, so we will isolate potential errors by starting the tasks' processes under a supervisor.

Adding a Supervisor

Elixir provides a *Supervisor behaviour* for creating supervisor processes, which you have to implement yourself. We are going to do that in Chapter 2, Long-Running Processes Using GenServer, on page 25. However, there are also some built-in supervisors which we can use without writing (almost) any code, and one of them is Task.Supervisor. It's made specifically for working with task processes, so it's an excellent choice for us. Open application.ex and update the children list:

```
sender/lib/sender/application.change1.ex
def start(_type, _args) do
  children = [
    {Task.Supervisor, name: Sender.EmailTaskSupervisor}
  ]

  opts = [strategy: :one_for_one, name: Sender.Supervisor]
  Supervisor.start_link(children, opts)
end
```

The element we added is referred to as a *child specification*. There are different types of formats for writing child specifications, but most of the times it is a tuple, containing information about the child process. This information is used by the supervisor to identify, start, and link the process.

Since Task.Supervisor is a built-in module, our child specification is simply the module name, followed by a list of options. We use the :name option to give it a unique name of our choice, which we will use later. Naming a process is known as *name registration*. It is common to append "Supervisor" when naming supervisors, so we called it Sender.EmailTaskSupervisor.

The child specification could be also be a map. We can re-write the last change and use a map like this:

```
def start(_type, _args) do
  children = [
    %{
      id: Sender.EmailTaskSupervisor,
      start: {
        Task.Supervisor,
        :start_link,
        [[name: Sender.EmailTaskSupervisor]]
      }
    }
  ]

  opts = [strategy: :one_for_one, name: Sender.Supervisor]
  Supervisor.start_link(children, opts)
end
```

Using a map is more verbose, but it allows you to set other configuration options in the map, in addition to the :id and :start values. There is also a helper function Supervisor.child_spec/1[2] which returns a map, and let's you override only the keys you need.

We're going to keep the tuple format for now, but we'll use the other formats too later in the book.

Our supervisor is now ready to use. It is as simple as that and there is no need to create any files or write more code. This simple change allows us to use a whole array of functions found in the Task.Supervisor module, some of which we introduced earlier.

---

2. https://hexdocs.pm/elixir/Supervisor.html#child_spec/2

Using Task.Supervisor

Before we switch to using EmailTaskSupervisor, let's see what happens when an error occurs and there is no supervisor in place. We are going to simulate an error by raising an exception when sending one of our fake emails. Edit sender.ex and add the following function clause just before send_email/1:

```
sender/lib/sender.change4.ex
def send_email("konnichiwa@world.com" = email), do:
  raise "Oops, couldn't send email to #{email}!"

def send_email(email) do
  Process.sleep(3000)
  IO.puts("Email to #{email} sent")
  {:ok, "email_sent"}
end
```

Using pattern matching, we are going to raise an exception only when the email address is konnichiwa@world.com. Let's see the impact of this error in practice. Restart your IEx shell and run the built-in self/0 function:

```
iex(1)> self()
#PID<0.136.0>
```

You get a process identifier for the current process, which is the IEx shell itself. Your number may be different than mine, but that's okay. Take a note of it.

Now, let's run our notify_all/1 function and see what happens:

```
iex(2)> Sender.notify_all(emails)

[error] Task #PID<0.154.0> started from #PID<0.136.0> terminating
** (RuntimeError) Oops, couldn't send email to konnichiwa@world.com!
  (sender) lib/sender.ex:8: Sender.send_email/1
  (elixir) lib/task/supervised.ex:90: Task.Supervised.invoke_mfa/2
  (elixir) lib/task/supervised.ex:35: Task.Supervised.reply/5
  (stdlib) proc_lib.erl:249: :proc_lib.init_p_do_apply/3
Function: &:erlang.apply/2
  Args: [#Function<0.104978293/1 in Sender.notify_all/1>,
  ["konnichiwa@world.com"]]
** (EXIT from #PID<0.136.0>) shell process exited with reason:
  an exception was raised:
  ** (RuntimeError) Oops, couldn't send email to konnichiwa@world.com!
    (sender) lib/sender.ex:8: Sender.send_email/1
    (elixir) lib/task/supervised.ex:90: Task.Supervised.invoke_mfa/2
    (elixir) lib/task/supervised.ex:35: Task.Supervised.reply/5
    (stdlib) proc_lib.erl:249: :proc_lib.init_p_do_apply/3

Interactive Elixir (1.9.4) - press Ctrl+C to exit (type h() ENTER for help)
iex(1)>
```

That's one big error message. In fact, there are two error messages that seem to be almost identical. Remember what we said about process links and how crashes propagate between linked processes? In this case, a task process linked to IEx raises an exception. Because of this link, created by async_stream/3, IEx also crashed with the same exception message. You can verify this by running self() again:

```
iex(1)> self()
#PID<0.155.0>
```

We have a new process identifier which means that a new IEx process was started and the old one has crashed.

We can prevent this by using the new EmailTaskSupervisor. Change notify_all/1 to use Task.Supervisor.async_stream_nolink/4 instead of Task.async_stream/3:

```
sender/lib/sender.change4.ex
def notify_all(emails) do
  Sender.EmailTaskSupervisor
  |> Task.Supervisor.async_stream_nolink(emails, &send_email/1)
  |> Enum.to_list()
end
```

The new code is very similar to the old one. The first argument of async_stream_nolink/4 is a supervisor module for the task processes—that's where you can use Sender.EmailTaskSupervisor. The rest of the arguments are identical to Task.async_stream/3. Let's run recompile() and notify_all/1 again:

```
iex> Sender.notify_all(emails)

[error] Task #PID<0.173.0> started from #PID<0.155.0> terminating
** (RuntimeError) Oops, couldn't send email to konnichiwa@world.com!
    (sender) lib/sender.ex:8: Sender.send_email/1
    (elixir) lib/task/supervised.ex:90: Task.Supervised.invoke_mfa/2
    (elixir) lib/task/supervised.ex:35: Task.Supervised.reply/5
    (stdlib) proc_lib.erl:249: :proc_lib.init_p_do_apply/3
Function: &:erlang.apply/2
    Args: [#Function<0.32410907/1 in Sender.notify_all/1>,
    ["konnichiwa@world.com"]]

[info]  Email to hello@world.com sent
[info]  Email to hola@world.com sent
[info]  Email to nihao@world.com sent
[
  ok: {:ok, "email_sent"},
  ok: {:ok, "email_sent"},
  ok: {:ok, "email_sent"},
  exit: {%RuntimeError{
    message: "Oops, couldn't send email to konnichiwa@world.com!"
  },
```

```
  [
    {Sender, :send_email, 1, [file: 'lib/sender.ex', line: 8]},
    {Task.Supervised, :invoke_mfa, 2,
     [file: 'lib/task/supervised.ex', line: 90]},
    {Task.Supervised, :reply, 5, [file: 'lib/task/supervised.ex', line: 35]},
    {:proc_lib, :init_p_do_apply, 3, [file: 'proc_lib.erl', line: 249]}
  ]}
]
```

The familiar exception is still printed out, but this time we got a meaningful result. We can see that all tasks have completed successfully, except one. The crashed task returned an :exit tuple containing the error message and even a stack trace. What's more important is that our current process was isolated from the crash—you can verify this by running self() again.

It's easy to see that this is a huge improvement over our previous version. The supervisor which we introduced handles concurrent processes and errors, while we're still reaping the performance benefits of async_stream.

You may have heard of a philosophy associated with Erlang called *let it crash*. Erlang applications can quickly recover from errors by restarting parts of its system. Elixir also embraces this philosophy. Unfortunately, this approach is also frequently misunderstood. In the next section we're going to explain what *let it crash* really means and how supervisors can restart processes to make our systems even more resilient.

Understanding Let It Crash

We used Task.Supervisor to isolate a process crash, but it may seem strange that we didn't prevent the crash by simply adding error handling in the send_email/1 function. In this particular case, we did this on purpose, just to simulate an unexpected exception. In practice, you should provide error handling when you expect an error to occur, and leave the rest to the supervisor as a last resort.

When discussing error handling for Elixir, the phrase *let it crash* is often used. As a result, some people assume that *let it crash* means that Erlang and Elixir developers don't do any error handling, which is not the case. Pattern matching and the with macro in Elixir make working with {:ok, result} and {:error, msg} tuples easy, and this approach is widely used in the community. Elixir also has try and rescue for catching exceptions, similarly to try and catch in other languages.

However, as much as we try as engineers, we know that errors can happen. This often leads to something called *defensive programming*. It describes the

practice of relentlessly trying to cover every single possible scenario for failure, even when some scenarios are very unlikely to happen, and not worth dealing with.

Erlang and Elixir take a different approach to defensive programming. Since all code runs in processes and processes that are lightweight, they focus on how the system can *recover from crashes* versus how to *prevent all crashes.* You can choose to allow a part (or even the whole) of the application to crash and restart, but handle other errors yourself. This shift in thinking and software design is the reason why Erlang became famous for its reliability and scalability.

That's where supervisors come into play. We saw that supervisors can isolate crashes, but they can also restart child processes. There are three different *restart values* available to us:

- :temporary option which never restarts processes;
- :transient will restart child processes, but only when they exit with an error;
- :permanent always restarts children, keeping them running, even when they try to shut down without an error;

> **Clarification on the terminology used**
>
> The word restart is slightly misleading in the context of Elixir processes. Once a process exits, it cannot be brought back to life. Therefore, restarting a process results in starting a new process to take the place of the old one, using the same child specification.

Depending on the frequency of the crash, the supervisor itself can also terminate when a child process cannot be recovered. Remember that supervisor's responsibility is to watch over its processes. If a :transient or :permanent restart value is used and a process keeps crashing, the supervisor will exit, because it has failed to restart that process. We are going to discuss this in more detail in Chapter 2, Long-Running Processes Using GenServer, on page 25.

The built-in Task.Supervisor which we used so far is already using the :temporary restart value for child processes. This is a sensible default, because it prevents all tasks processes from crashing if another task exits with an error. However, these options will come in handy in the next chapter when we start building more complex supervision trees.

Wrapping Up

Congratulations on completing this chapter! You've come a long way from running synchronous code and patiently waiting for it to complete. We covered

Elixir processes and used the Task module to do work concurrently and in parallel, easily avoiding timeout errors and unexpected exceptions. You also learned about supervisors, which are the foundation of building fault-tolerant Elixir applications. A lot of the things we covered so far will appear again in the following chapters, so we have created a strong foundation to built upon.

Although the Task module is very powerful and versatile, it is only useful for running one-off functions concurrently. As a result, Task processes are short-lived and exit as soon as they complete. In the next chapter, we're going to introduce another type of process, which is capable of running and maintaining state for as long as we need it to. It is going to completely change the way you build applications forever, and this is not an exaggeration. We're also going to expand our knowledge of supervisors and start building more complex systems. Let's not waste time and jump straight into it.

Long-Running Processes Using GenServer

The Task module is useful for running single async functions, but as logic becomes more complex, you will need a sharper tool. In this chapter, we'll look at how to create long-lived processes that also run in the background, but offer greater control and more flexibility.

Concurrent work often takes a long time to complete, such as when you are importing a large amount of user data from another service or relying on third-party APIs, for example. This presents several challenges when using the Task module. What if an API service goes briefly offline or a network error causes the task to fail? Do you need to increase the :timeout setting because the network is slow? How do you show progress to the user and provide better user experience?

There are also times when we want to run some code in the background continuously. In a banking application, you may need the latest foreign currency exchange rates, which means periodically retrieving the data and storing it somewhere, for as long as the application runs.

These and many other problems could be solved by using a GenServer. GenServer, which is short for *generic server*, enables us to create concurrent processes that we can interact with. Just like a Web server, it listens for requests and can respond with a result. GenServer processes also have their own state, which they keep in-memory until the process exits. You also have a lot more options when it comes to supervising GenServer processes and dealing with potential errors. Since GenServer is part of the OTP, it is already available for you to use in Elixir—there is no need to install any dependencies.

In this chapter, you are going to learn how to use GenServer to create stateful processes that you can interact with. We'll look at the different supervisor strategies and see how restart options work in practice. We'll tie everything

together by building a simple but effective job processing system, and you'll understand how to design fault-tolerant applications and manage processes using the Elixir Registry.

There is a lot to cover in this chapter, so let's get started!

Starting With a Basic GenServer

We are going to continue working on the sender project from the previous chapter. If you decided to skip that part, please see Creating Our Playground, on page 3 and follow the steps to create a new Elixir project and make the required change in sender.ex. We are going to create a GenServer process from scratch and see how it compares to Task processes.

First, we're going to create a new file in the lib directory. Let's name it send_server.ex. In that file, we'll define our SendServer module like so:

```
defmodule SendServer do
  use GenServer

end
```

The use macro for the GenServer module does two things for us. First, it automatically injects the line @behaviour GenServer in our SendServer module. If you're not familiar with behaviours in Elixir, they're similar to interfaces and contracts in other programming languages. Second, it provides a default GenServer implementation for us by injecting all functions required by the GenServer behaviour.

We can verify that our code compiles by starting IEx:

```
$ iex -S mix
Compiling 3 files (.ex)
warning: function init/1 required by behaviour GenServer is not
  implemented (in module SendServer).

We will inject a default implementation for now:

    def init(init_arg) do
      {:ok, init_arg}
    end

You can copy the implementation above or define your own that converts the
  arguments given to GenServer.start_link/3 to the server state.

  lib/send_server.ex:1: SendServer (module)

Generated sender app
Interactive Elixir (1.9.4) - press Ctrl+C to exit (type h() ENTER for help)
iex(1)>
```

You just created your first GenServer process! However, we didn't write any logic for SendServer so it doesn't do anything useful at this point. We also got a warning message by the compiler, telling us that the init/1 function is missing and was replaced by a default implementation.

The GenServer module provides default implementations for several functions required by the GenServer behaviour. These functions are known as *callbacks*, and init/1 is one of them. Callbacks are important because they allow you to customize the GenServer process by adding your own business logic to it. In the following section, you will learn about the most frequently used callbacks and some examples of how to use them.

GenServer Callbacks In-Depth

The best way to learn how callbacks work is to see them in action. We are going to add some functionality to the SendServer module and introduce the most common GenServer callbacks along the way.

You can implement a callback by declaring it in the SendServer module like any other function. By doing this, you are replacing the default function that GenServer provides. This is sometimes called *overriding* the default implementation. When implementing a callback, there are two things you need to know:

- what arguments the callback function takes
- what return values are supported

We are going to cover the following callback functions for the GenServer behaviour:

- handle_call/3
- handle_cast/2
- handle_continue/2
- handle_info/2
- init/1
- terminate/2

Learning about these callbacks will enable you take full advantage of your GenServer process. There are two callbacks which we're going to skip, code_change/2 and format_status/2. We will also focus on the most frequently used return values, leaving some behind. All callbacks and return values are well documented on Elixir's HexDocs page, so feel free to go online and explore further what's possible with GenServer.

Exploring callbacks for behaviours

The easiest way to find out what callbacks are available for a specific behaviour is to check the documentation for the module on HexDocs. For example, you can go to Elixir's own HexDocs page.[1] Using the left sidebar, navigate to Modules and find the module name in the list of categories. GenServer can be found under "Processes & Applications". When you click on the module name, it will expand and show you a few sub-items: Top, Sections, Summary, Types, Functions, and Callbacks.

Click on Callbacks and you will see the links of the functions available for you to implement.

Remember the warning about the init/1 function that we saw earlier? Let's fix it by implementing the init/1 callback first.

Initializing the Process

The init/1 callback runs as soon as the process starts. When you start a GenServer process, you can optionally provide a list of arguments. This list is made available to you in the init/1 callback. This is a convenient way to provide some configuration details at runtime for the process. We mentioned that each GenServer process has its own in-memory state. This state is created by the init/1 function as well.

We are going to extend SendServer with the ability to send emails. If an email fails to send for whatever reason, we will also retry sending it, but we will limit the maximum number of retries. We will keep the business logic as simple as possible, so we can focus on learning how GenServer works.

First, let's add the following code to send_server.ex right after use GenServer:

```elixir
def init(args) do
  IO.puts("Received arguments: #{inspect(args)}")
  max_retries = Keyword.get(args, :max_retries, 5)
  state = %{emails: [], max_retries: max_retries}
  {:ok, state}
end
```

We use the argument max_retries if present, otherwise we default to 5 retries max. We will also keep track of all sent emails using the emails list. These variables will be kept in the initial state for the process. Finally, the function returns {:ok, state}. This means that the process has successfully initialized.

---

1. https://hexdocs.pm/elixir

We will cover the other possible return values in just a moment, but first, let's start SendServer to make sure it works as expected.

With the IEx shell open, run recompile() and then the following code:

```
iex> {:ok, pid} = GenServer.start(SendServer, [max_retries: 1])
```

You should see output similar to this one:

```
Received arguments: [max_retries: 1]
{:ok, #PID<0.228.0>}
```

SendServer is now running in the background. There are several ways to stop a running process, as you will see later, but for now, let's use GenServer.stop/3:

```
iex> GenServer.stop(pid)
:ok
```

There are a number of result values supported by the init/1 callback. The most common ones are:

```
{:ok, state}
{:ok, state, {:continue, term}}
:ignore
{:stop, reason}
```

We already used {:ok, state}. The extra option {:continue, term} is great for doing post-initialization work. You may be tempted to add complex logic to your init/1 function, such as fetching information from the database to populate the GenServer state, but that's not desirable because the init/1 function is synchronous and should be quick. This is where {:continue, term} becomes really useful. If you return {:ok, state, {:continue, :fetch_from_database}}, the handle_continue/2 callback will be invoked after init/1, so you can provide the following implementation:

```
def handle_continue(:fetch_from_database, state) do
  # called after init/1
end
```

We will discuss handle_continue/2 in just a moment.

Finally, the last two return values help us stop the process from starting. If the given configuration is not valid or something else prevents this process from continuing, we can return either :ignore or {:stop, reason}. The difference is that if the process is under a supervisor, {:stop, reason} will make the supervisor restart it, while :ignore won't trigger a restart.

Breaking Down Work in Multiple Steps

The handle_continue/2 callback is a recent addition to GenServer. Often GenServer processes do complex work as soon as the they start. Rather than blocking the whole application from starting, we return {:ok, state, {:continue, term}} from the init/1 callback, and use handle_continue/2.

Accepted return values for handle_continue/2 include:

```
{:noreply, new_state}
{:noreply, new_state, {:continue, term}}
{:stop, reason, new_state}
```

Since the callback receives the latest state, we can use it to update it with new information by returning {:noreply, new_state}. For example:

```
def handle_continue(:fetch_from_database, state) do
  # get `users` from the database
  {:noreply, Map.put(state, :users, users)}
end
```

The other return values are similar to the ones we already covered for init/1, but it's interesting to note that handle_continue/2 can also return {:continue, term}, which will trigger another handle_continue/2. You can use this to break down work to several steps when needed. Although handle_continue/2 is often used in conjunction with init/1, other callbacks can also return {:continue, term}.

Next, we're going to see how we can communicate with the GenServer process.

Sending Process Messages

One of the highlights of GenServer processes is that you can interact with them while they're still running. This is done by sending messages to the process. If you want to get some information back from the process, you use GenServer.call/3. When you don't need a result back, you can use GenServer.cast/2. Both functions accept the process identifier as their first argument, and a message to send to the process as the second argument. Messages could be any Elixir term.

When the cast/2 and call/3 functions are used, the handle_cast/2 and handle_call/3 callbacks are invoked, respectively. Let's see how they work in practice. We will implement an interface to get the current process state, as well as start sending emails.

Add the following code to send_server.ex:

```
def handle_call(:get_state, _from, state) do
  {:reply, state, state}
```

```
end
```

It is common to use pattern matching when implementing callbacks, since there could be multiple callback implementations for each type of message. For this one, we expect the message :get_state. The arguments given to handle_call/3 include the sender (which we do not use, hence the underscore _from) and the current process state.

The most common return values from handle_call/3 are:

```
{:reply, reply, new_state}
{:reply, reply, new_state, {:continue, term}}
{:stop, reason, reply, new_state}
```

By returning {reply, state, state} we send back the current state to the caller. Let's recompile() and try it in action:

```
iex> {:ok, pid} = GenServer.start(SendServer, [max_retries: 1])
Received arguments: [limit: 1]
{:ok, #PID<0.265.0>}

iex> GenServer.call(pid, :get_state)
%{emails: [], max_retries: 1}
```

That's great, we can see that the current state is just as we set it up in the init/1 callback. Now, let's implement sending emails using handle_cast/2. The arguments given to handle_cast/2 are just a term for the message and state. We're going to pattern match on the message {:send, email}:

```
def handle_cast({:send, email}, state) do
  # to do...
end
```

Most of the times you will return one of the following tuples:

```
{:noreply, new_state}
{:noreply, new_state, {:continue, term}}
{:stop, reason, new_state}
```

To refresh your memory, this is how the Sender.send_email/1 function works:

```
def send_email(email) do
  Process.sleep(3000)
  IO.puts("Email to #{email} sent")
  {:ok, "email_sent"}
end
```

Let's call the send_email/1 function from our handle_cast/2 callback, and update the process state when this happens:

```
def handle_cast({:send, email}, state) do
  Sender.send_email(email)
```

```elixir
  emails = [%{email: email, status: "sent", retries: 0}] ++ state.emails

  {:noreply, state}
end
```

To try out these changes, you may need to restart IEx and run the process again. Let's see if sending emails works:

```elixir
iex(1)> {:ok, pid} = GenServer.start(SendServer, [max_retries: 1])
Received arguments: [max_retries: 1]
{:ok, #PID<0.151.0>}

iex(2)> GenServer.cast(pid, {:send, "hello@email.com"})
:ok
Email to hello@email.com sent

iex(3)> GenServer.call(pid, :get_state)
%{
  emails: [%{email: "hello@email.com", retries: 0, status: "sent"}],
  max_retries: 1
}
```

Everything seems to work just as we wanted. When using GenServer.cast/2, we always get :ok as a reply. In fact, the reply comes almost immediately. This simply means that the GenServer process has acknowledged the message, while the actual work is being performed by the process.

Notifying the Process of Events

Other than using GenServer.cast/2 and GenServer.call/3, you can also send a message to a process using Process.send/2. This type of generic message will trigger the handle_info/2 callback, which works exactly like handle_cast/2, and can return the same set of tuples. Usually handle_info/2 deals with system messages. Normally, you will expose your server API using cast/2 and call/2, and keep send/2 for internal use.

Let's see how we can use handle_info/2. We will implement retries for emails that fail to send. To test this, we need to modify sender.ex, so one of the emails returns an error. Replace your send_mail/1 logic with this:

```elixir
def send_email("konnichiwa@world.com" = email),
  do: :error

def send_email(email) do
  Process.sleep(3000)
  IO.puts("Email to #{email} sent")
  {:ok, "email_sent"}
end
```

Now all emails to konnichiwa@world.com will return :error, but we have to make sure this is persisted correctly. Next, update the handle_cast/2 callback in Send-Server:

```elixir
def handle_cast({:send, email}, state) do
  status =
    case Sender.send_email(email) do
      {:ok, "email_sent"} -> "sent"
      :error -> "failed"
    end

  emails = [%{email: email, status: status, retries: 0}] ++ state.emails

  {:noreply, Map.put(state, :emails, emails)}
end
```

Now we have everything in place to implement retries. We will use Process.send_after/3 which is similar to Process.send/2, except that it sends the message after the specified delay. We will start periodically checking for failed emails as soon as the server starts, so let's add this to our init/1 callback, before the return statement:

```elixir
Process.send_after(self(), :retry, 5000)
```

The first argument of send_after/3 is the process identifier. We use the self() function which returns the *PID* for the current process. The second argument is the name of the message, which in our case is :retry, but it could be any Elixir term. The last argument is the delay, in milliseconds, after which the message will be sent. We set this to five seconds.

Finally, let's implement handle_info/2 for the :retry message. This implementation will be a bit bigger in comparison to previous ones, but don't worry, we'll explain everything right after:

```elixir
def handle_info(:retry, state) do
  {failed, done} =
    Enum.split_with(state.emails, fn item ->
      item.status == "failed" && item.retries < state.max_retries
    end)

  retried =
    Enum.map(failed, fn item ->
      IO.puts("Retrying email #{item.email}...")

      new_status =
        case Sender.send_email(item.email) do
          {:ok, "email_sent"} -> "sent"
          :error -> "failed"
        end

      %{email: item.email, status: new_status, retries: item.retries + 1}
```

```
    end)
  Process.send_after(self(), :retry, 5000)
  {:noreply, Map.put(state, :emails, retried ++ done)}
end
```

We use Enum.split_with/2 to divide the items in the state.emails list in two categories: emails to retry, and everything else. We don't want to retry emails that were sent successfully, or the ones that exceeded the max_retries setting for SendServer. The rest of the code simply iterates over the failed group of emails and attempts to re-send them. We increase the retries counter after each attempt.

Finally, before we update the return and update the state, we schedule another periodic check with Process.send_after/3, so we can repeat the process.

You can now try this in IEx, just remember to use konnichiwa@world.com to test the new retry logic:

```
iex(1)> {:ok, pid} = GenServer.start(SendServer, max_retries: 2)
Received arguments: [max_retries: 2]
{:ok, #PID<0.156.0>}

iex(2)> GenServer.cast(pid, {:send, "hello@world.com"})
:ok
Email to hello@world.com sent

iex(3)> GenServer.cast(pid, {:send, "aloha@world.com"})
:ok
Email to aloha@world.com sent

iex(4)> GenServer.cast(pid, {:send, "konnichiwa@world.com"})
:ok
Retrying email konnichiwa@world.com...
Retrying email konnichiwa@world.com...

iex(5)> GenServer.call(pid, :get_state)
%{
  emails: [
    %{email: "konnichiwa@world.com", retries: 2, status: "failed"},
    %{email: "aloha@world.com", retries: 0, status: "sent"},
    %{email: "hello@world.com", retries: 0, status: "sent"}
  ],
  max_retries: 2
}
```

Great, this seems to work as expected. The process retries konnichiwa@world.com twice, before giving up. Note that when we started the process, we configured it with max_retries: 2 rather than falling back to the default value of 5. We can verify this by getting the process' state using GenServer.call/2.

Process Teardown

The final callback in the list is terminate/2. It is usually invoked before the process exits, but only when the process *itself* is responsible for the exit. Most often, the exit will result from explicitly returning {:stop, reason, state} from a callback (excluding init/1), or when an unhandled exception happens within the process.

Caution when relying on terminate/2

There are cases when a GenServer process is forced to exit due to an external event, for example, when the whole application shuts down. In those cases, terminate/2 will not be invoked by default. This is something you have to keep in mind if you want to ensure that important business logic always runs before the process exits. For example, you may want to persist in-memory data (stored in the process' state) to the database before the process dies.

This is out of the scope of this book, but if you want to ensure that terminate/2 is always called, look into setting Process.flag(:trap_exit, true) on the process, or use Process.monitor/1 to perform the required work in a separate process.

Let's implement terminate/2 by printing a message before the process exits:

```
def terminate(reason, _state) do
  IO.puts("Terminating with reason #{reason}")
end
```

The return value of terminate/2 is not important, you can return anything. Restart IEx and use GenServer.stop/1 to stop the process:

```
iex(1)> {:ok, pid} = GenServer.start(SendServer, [])
Received arguments: []
{:ok, #PID<0.152.0>}

iex(2)> GenServer.stop(pid)
Terminating with reason normal
:ok
```

We covered the most popular callbacks and their return values, which gives you a strong foundation to start building your own GenServer processes. We also finished building our SendServer module, which can now send and retry failed emails.

However, our SendServer implementation is far from ideal. Remember that send_email/1 pauses the process for three seconds. If you try running GenServer.cast(pid, {:send, "hello@email.com"}) several times in a row, GenServer will

acknowledge all messages by returning :ok. However, it won't be able to perform the work straight away, if it is busy doing something else. After all, it is just a single process doing all the work. All these messages will put in the process' *message queue*, and executed one by one.

Even worse, if you try using GenServer.call(pid, :get_state) while the process is still busy, you will get a timeout error. By default, the GenServer.call/3 will error after 5000ms. We can tweak this by providing an optional third argument, but this is still not a practical solution.

We are going to fix this right away by building a significantly improved job processing system. It will be capable of performing any job concurrently, such as sending emails and retrying the jobs that fail. All we have to do is change our approach slightly. Keep reading!

Using the Task module with GenServer

You can free up your GenServer process by using Task.start/1 and Task.async/2 to run code concurrently. Upon completion, the Task process will send messages back to GenServer, which you can process with handle_info/2. For more information, see the documentation for Task.Supervisor.async_nolink/3[2] which contains some useful examples.

Building a Job Processing System

In the previous chapter, we used the Task module to start a new process when sending each email. We are going to adopt the same approach for our job processing system. Rather than having a single GenServer that does all the work, we will start a GenServer process for each job. This will help us leverage concurrency and improve performance.

We also don't want to lose any features that we have introduced so far in this chapter—we want to be able to retry failed jobs and configure the process to suit our needs. Finally, we will introduce a new type of supervisor to help us manage these processes and provide fault-tolerance. This figure will give you a high-level overview of what we're going to build:

---

2. https://hexdocs.pm/elixir/Task.Supervisor.html#async_nolink/3

```
                    ╭─────────────────╮
                    │   Application   │
                    ╰─────────────────╯
                             │
                    ╭─────────────────╮
                    │   Job runner    │
                    ╰─────────────────╯
                   ╱         │         ╲
            ╭──────────╮ ╭──────────╮ ╭──────────╮
            │  Job*    │ │  Job*    │ │  Job*    │
            ╰──────────╯ ╰──────────╯ ╰──────────╯

                 * processes started on demand
```

To begin, let's scaffold a new Elixir project with a supervision tree. We will call it jobber, which is not the most creative name, but it will do the job (pun intended):

```
$ mix new jobber --sup
```

Then change your current directory to jobber.

Next, we will create our GenServer process, which will do the heavy lifting. Create the file job.ex in lib/jobber with the following content:

```
defmodule Jobber.Job do
  use GenServer
  require Logger

end
```

Just like we did before, we brought in the GenServer behaviour and implementation. We are also going to use Elixir's Logger module for nicer log messages.

Next, let's implement the init/1 callback and initial process configuration.

Initializing the Job Process

We are going to use a struct to keep track of the process state. Structs are just like maps, but they provide some additional features, including default values for the struct's attributes. Each job will have several fields, such as a short id to uniquely identify the process, the work function itself, and the

already familiar retries and max_retries. We will also keep track of the status of each job, which will be either new, errored, failed or done.

Add the following struct definition to job.ex:

jobber/lib/jobber/job.ex
```elixir
defstruct [:work, :id, :max_retries, retries: 0, status: "new"]
```

Next, implement the init/1 callback:

jobber/lib/jobber/job.ex
```elixir
def init(args) do
  work = Keyword.fetch!(args, :work)
  id = Keyword.get(args, :id, random_job_id())
  max_retries = Keyword.get(args, :max_retries, 3)

  state = %Jobber.Job{id: id, work: work, max_retries: max_retries}
  {:ok, state, {:continue, :run}}
end
```

The code is similar to SendServer. You may notice that we use the Keyword.fetch!/2 function here to get the work argument. This function will throw an error if you forget to provide this parameter by accident. We also generate a default id, in case an id is not available. There are different ways to generate a random identifier in Elixir. We will use Erlang's :crypto module and define random_job_id/0 like this:

jobber/lib/jobber/job.ex
```elixir
defp random_job_id() do
  :crypto.strong_rand_bytes(5) |> Base.url_encode64(padding: false)
end
```

This will ensure that the id is unique and short at the same time, making it easy to work with.

Performing Work

We returned {:ok, state, {:continue, :run}} from the init/1 callback, so we have to handle this message next. Within the handle_continue/2 callback, we will perform the required work as soon as the process finishes initializing:

jobber/lib/jobber/job.ex
```elixir
def handle_continue(:run, state) do
  new_state = state.work.() |> handle_job_result(state)

  if new_state.status == "errored" do
    Process.send_after(self(), :retry, 5000)
    {:noreply, new_state}
  else
    Logger.info("Job exiting #{state.id}")
    {:stop, :normal, new_state}
```

```
      end
end
```

Since the work attribute is an anonymous function, we can invoke it using the state.work.() syntax, and then process the result. We are going to use the status attribute to determine whether or not we need to retry the job again. If the status is errored, we schedule another attempt in five seconds, and keep the process running. Otherwise, for done and failed, we stop the process, as there isn't any work left to do.

We need to considered three different outcomes when handling the result:

- success, when the job completes and returns {:ok, data}
- initial error, when it fails the first time with :error
- retry error, when we attempt to re-run the job and also receive :error

We will use pattern matching here, so each handle_job_result/2 function definition will handle one outcome:

```
jobber/lib/jobber/job.ex
defp handle_job_result({:ok, _data}, state) do
  Logger.info("Job completed #{state.id}")
  %Jobber.Job{state | status: "done"}
end

defp handle_job_result(:error, %{status: "new"} = state) do
  Logger.warn("Job errored #{state.id}")
  %Jobber.Job{state | status: "errored"}
end

defp handle_job_result(:error, %{status: "errored"} = state) do
  Logger.warn("Job retry failed #{state.id}")
  new_state = %Jobber.Job{state | retries: state.retries + 1}

  if new_state.retries == state.max_retries do
    %Jobber.Job{new_state | status: "failed"}
  else
    new_state
  end
end
```

The logic is again not much different compared to SendServer. However, this implementation requires that all jobs return {:ok, data} (although we ignore the data result to keep things simple) or :error. If there is more than one type of error that you want support, you can return {:error, reason} instead, to allow for more granular error handling.

Finally, we have to implement a callback to process the :retry message. Since retrying a task is the same as running it again, we can do this:

```
jobber/lib/jobber/job.ex
def handle_info(:retry, state) do
  # Delegate work to the `handle_continue/2` callback.
  {:noreply, state, {:continue, :run}}
end
```

This will immediately trigger another :run message.

We're now ready to see how the Job process works in action. Start IEx and run a job like this:

```
iex(1)> GenServer.start(Jobber.Job, work: fn -> Process.sleep(5000);
  {:ok, []} end)

{:ok, #PID<0.261.0>}

11:35:08.221 [info]  Job completed r4Fl3I4
11:35:08.221 [info]  Job exiting r4Fl3I4
```

We used Process.sleep/1 to pretend we're doing some time-consuming work, and returned {:ok, []} as a result. Notice the semicolon before the result. We're creating an anonymous function on a single line, so we need to separate the two statements, otherwise we will get a syntax error. You can create the anonymous function beforehand to make things easier:

```
iex(2)> good_job = fn ->
...(2)> Process.sleep(5000)
...(2)> {:ok, []}
...(2)> end
#Function<21.126501267/0 in :erl_eval.expr/5>
```

You can press Enter after fn -> to go on a new line in IEx. Now you can pass the good_job variable for the work argument directly:

```
GenServer.start(Jobber.Job, work: good_job)
```

To test the retry logic, we need a function that returns an :error. Let's create another mock function:

```
iex(3)> bad_job = fn ->
...(3)> Process.sleep(5000)
...(3)> :error
...(3)> end
#Function<21.126501267/0 in :erl_eval.expr/5>

iex(4)> GenServer.start(Jobber.Job, work: bad_job)
{:ok, #PID<0.263.0>}

11:38:49.218 [warn]  Job errored JJWXUVE
11:38:59.319 [warn]  Job retry failed JJWXUVE
11:39:09.420 [warn]  Job retry failed JJWXUVE
11:39:19.521 [warn]  Job retry failed JJWXUVE
11:39:19.521 [info]  Job exiting JJWXUVE
```

Everything is working as expected, which is good news. This is already a great improvement over SendServer, but we are not finished yet. Did you notice how we used GenServer.start/2 to start our GenServer process? This is similar to Task.start/1 in a way that it starts a process, without linking it to the parent process. However, if a Job process crashes due to an error, we're going to lose the process.

We can opt for using GenServer.start_link/2 instead. But now we have a different problem. The parent process will be linked, and any crash from a Job process will also propagate to the parent.

The best solution is to start all Job processes under a supervisor. The supervisor will take care of restarting the processes in the event of an error, and also start jobs on demand. To accomplish this, we are going to use a new type of supervisor.

Introducing DynamicSupervisor

DynamicSupervisor is another ready-to-use supervisor available to you. It can start any GenServer process on demand. Let's add it to the supervision tree:

jobber/lib/jobber/application.ex
```
def start(_type, _args) do
  children = [
    {DynamicSupervisor, strategy: :one_for_one, name: Jobber.JobRunner},
  ]

  opts = [strategy: :one_for_one, name: Jobber.Supervisor]
  Supervisor.start_link(children, opts)
end
```

The strategy setting is required, and the only strategy that is currently accepted is :one_for_one. We will talk about supervisor strategies later in this chapter, so don't worry about this for now.

> **Module-Based DynamicSupervisor**
>
> You can also define a DynamicSupervisor module, like we do later in Implementing a Supervisor, on page 45. Instead of use Supervisor you will need use DynamicSupervisor and call DynamicSupervisor.init/1 with the required :strategy value. You don't have to provide a list of children processes.

It will be useful to create some helpers for testing. Create .iex.exs at the top project directory with the following content:

jobber/.iex.exs
```
good_job = fn ->
```

```
    Process.sleep(5000)
    {:ok, []}
end

bad_job = fn ->
  Process.sleep(5000)
  :error
end
```

Pre-defining the test functions will save us some typing.

Next, we're going to create a helper function for starting jobs in jobber.ex:

jobber/lib/jobber.ex
```
defmodule Jobber do
  alias Jobber.{JobRunner, Job}

  def start_job(args) do
    DynamicSupervisor.start_child(JobRunner, {Job, args})
  end
end
```

Let's restart IEx and make sure everything works:

```
iex(1)> Jobber.start_job(work: good_job)
{:error,
 {:undef,
  [
    {Jobber.Job, :start_link,
     [[work: #Function<21.126501267/0 in :erl_eval.expr/5>]], []},
    {DynamicSupervisor, :start_child, 3,
     [file: 'lib/dynamic_supervisor.ex', line: 690]},
    {DynamicSupervisor, :handle_start_child, 2,
     [file: 'lib/dynamic_supervisor.ex', line: 676]},
    {:gen_server, :try_handle_call, 4, [file: 'gen_server.erl', line: 661]},
    {:gen_server, :handle_msg, 6, [file: 'gen_server.erl', line: 690]},
    {:proc_lib, :init_p_do_apply, 3, [file: 'proc_lib.erl', line: 249]}
  ]}}
```

DynamicSupervisor.start_child/2 was supposed to start a new Job process, but it failed with an error. The error message is not very clear, but the reason is that we haven't implemented a start_link/1 function in the Job module yet. DynamicSupervisor expects start_link/1 to be defined, and uses it to start the process and link it automatically. Let's add it next:

jobber/lib/jobber/job.ex
```
def start_link(args) do
  GenServer.start_link(__MODULE__, args)
end
```

This simple fix should be enough. Run recompile() in IEx and try it again. This time, the job will start successfully, but then something strange happens:

```
iex(1)> Jobber.start_job(work: good_job)
{:ok, #PID<0.163.0>}

17:02:28.773 [info]  Job completed nBSU4U0
17:02:28.780 [info]  Job exiting nBSU4U0
17:02:33.782 [info]  Job completed PkDVtBs
17:02:33.782 [info]  Job exiting PkDVtBs
17:02:38.783 [info]  Job completed Q0MS8T8
...
```

You get an endless stream of jobs being started. What could be the issue?

Process Restart Values Re-Visited

In the previous chapter in Understanding Let It Crash, on page 22 we talked about restart values. By default, GenServer processes are always restarted by their supervisor, which is the :permanent setting. In our case, we intentionally shut down the process. However, JobRunner thinks something must have gone wrong, so it will keep restarting the process forever. We can easily fix this by using the :transient restart option, which tells the supervisor not to restart the process if it is exiting normally. Modify our use GenServer statement in job.ex to set the new setting:

```
use GenServer, restart: :transient
```

You can quit the old IEx session and start a new one to give this a go:

```
iex(1)> Jobber.start_job(work: good_job)
{:ok, #PID<0.158.0>}

17:28:41.456 [info]  Job completed fJ1qlbo
17:28:41.456 [info]  Job exiting fJ1qlbo
```

We fixed the issue successfully. Let's try the bad_job test case:

```
iex(2)> Jobber.start_job(work: bad_job)
{:ok, #PID<0.160.0>}

14:31:14.798 [warn]  Job errored DFXanZw
14:31:24.899 [warn]  Job retry failed DFXanZw
14:31:35.000 [warn]  Job retry failed DFXanZw
14:31:45.101 [warn]  Job retry failed DFXanZw
14:31:45.101 [info]  Job exiting DFXanZw
```

So far so good, but the ultimate test is to see how the JobRunner supervisor performs when an exception happens. Exit IEx and edit .iex.exs to add a new test case function:

```
doomed_job = fn ->
  Process.sleep(5000)
  raise "Boom!"
end
```

Start the jobber app again and give the doomed_job function to JobRunner:

```
iex> Jobber.start_job(work: doomed_job)
{:ok, #PID<0.163.0>}

[error] GenServer #PID<0.163.0> terminating
** (RuntimeError) Boom!
    (stdlib 3.13) erl_eval.erl:678: :erl_eval.do_apply/6
    (jobber 0.1.0) lib/jobber/job.ex:28: Jobber.Job.handle_continue/2
    (stdlib 3.13) gen_server.erl:680: :gen_server.try_dispatch/4
    (stdlib 3.13) gen_server.erl:431: :gen_server.loop/7
    (stdlib 3.13) proc_lib.erl:226: :proc_lib.init_p_do_apply/3
Last message: {:continue, :run}

-- rest of the error log omitted --
```

The process is stuck once gain. It seems to be continuously restarting by throwing the same exception over and over again. There is one setting that we overlooked—the restart configuration on the supervisor. We will tackle this next.

Adjusting Restart Frequency

A supervisor like JobRunner has two settings—max_restarts and max_seconds—that control something we call *restart frequency*. By default, max_restarts is set to 3 and max_seconds to 5. As a result, during a five second time window, a supervisor will attempt to restart a process three times, before it gives up.

Right now, our Process.sleep/1 call in doomed_job is set to stop the process after 5000ms. This means that the supervisor will have time for only one restart, until the five seconds window expires. Since our fake function is taking a fixed amount of time to run and error, this process of restarting continues forever.

Let's increase max_seconds to thirty seconds. Go to application.ex and add the extra setting:

jobber/lib/jobber/application.change1.ex
```elixir
def start(_type, _args) do
  job_runner_config = [
    strategy: :one_for_one,
    max_seconds: 30_000,
    name: Jobber.JobRunner
  ]

  children = [
    {DynamicSupervisor, job_runner_config}
  ]

  opts = [strategy: :one_for_one, name: Jobber.Supervisor]
  Supervisor.start_link(children, opts)
```

```
end
```

Start IEx again, but this time, before you try running the experiment again, let's take a note of the PID of the JobRunner process. You will see why in a second. You can use Process.whereis/1 to retrieve the process identifier using the name of the process:

```
iex(1)> Process.whereis(Jobber.JobRunner)
#PID<0.141.0>
```

Let's try running our faulty function again. You should see something like this:

```
iex(2)> Jobber.start_job(work: doomed_job)
{:ok, #PID<0.146.0>}

16:51:59.814 [error] GenServer #PID<0.146.0> terminating
** (RuntimeError) Boom!

-- rest of the error log omitted --
```

You will see the error message above four times, because the process is restarted three times after the first exception. This may look like a success, but it's too early to celebrate. Let's check the PID for JobRunner again:

```
iex(3)> Process.whereis(Jobber.JobRunner)
#PID<0.152.0>
```

As you can see, there is a new process identifier for JobRunner, which means that the supervisor itself was restarted. Remember that supervisor's responsibility is to watch over its processes. If a :transient or :permanent restart option is used and a process keeps crashing, the supervisor will exit, because it has failed to restart that process and ensure the reliability of the system.

As a result of JobRunner exiting, any concurrent Job processes will be also terminated. Even if a single process fails to be recovered, it could endanger all other running processes. However, this can be easily fixed by adding a supervisor for each Job process, which will handle restarts and gracefully exit when its process fails.
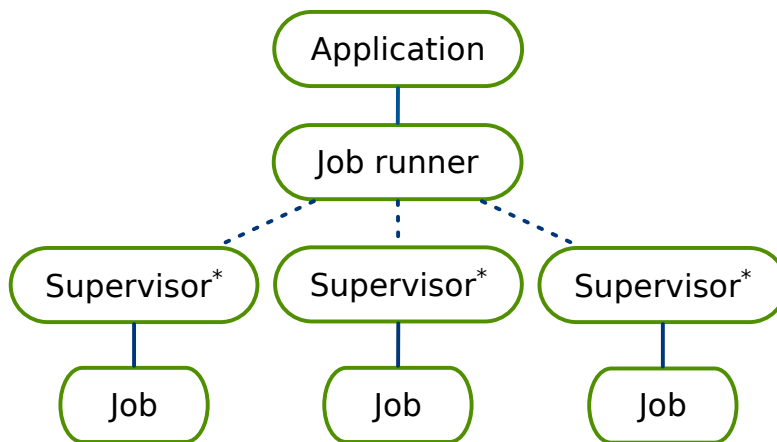
Implementing a Supervisor

So far we have used two types of supervisors—Task.Supervisor and DynamicSupervisor. They're great because they allow you to start processes on demand. However, they are both made with a specific goal in mind, and sometimes you will need greater configurability over how the supervisor works.

This is why you will often reach to the Supervisor behaviour to create your own supervisor from scratch. In this section, we will create a supervisor using the Supervisor behaviour, and put it in front of each Job process.

We already learned that supervisors can restart child processes, and we saw how the :permanent and :transient restart values work. The last setting we haven't used is :temporary. This value tells the supervisor to ignore the process and don't worry about restarting it, even when the child process exits with an error.

Our goal is to make JobRunner start an intermediary supervisor for each job, which in turn will start the actual Job process. This new supervisor will restart the Job process if needed, but will have a restart value of :temporary, so it won't cause further damage if it fails. This figure shows what the supervision tree is going to look like after we make all the changes:

* processes started on demand

Let's start by creating the supervisor module:

```elixir
defmodule Jobber.JobSupervisor do
  use Supervisor, restart: :temporary

  # TODO
end
```

The use Supervisor macro is similar to use GenServer. It declares that we are implementing the Supervisor behaviour. We also make sure that this process

has the :temporary restart value. Since this is going to be a linked process, we'll implement start_link/1 and init/1 next:

jobber/lib/jobber/job_supervisor.ex
```
def start_link(args) do
  Supervisor.start_link(__MODULE__, args)
end

def init(args) do
  children = [
    {Jobber.Job, args}
  ]

  options = [
    strategy: :one_for_one,
    max_seconds: 30_000
  ]

  Supervisor.init(children, options)
end
```

These functions should look familiar to what we do when creating a GenServer, but there are some differences as well. The JobSupervisor process also takes arguments, which are handed over to the init/1 function. Then we define a list of children, which for JobSupervisor is just a single Job process. In this case, we want to pass down all configuration to the child process. Finally, instead of returning a tuple, we delegate the rest of the work to Supervisor.init/2, which will start the supervisor. If you open application.ex, you will see that the main application supervisor works the same way.

Finally, let's modify Jobber.start_job/1 to use JobSupervisor instead of running the Job process directly:

jobber/lib/jobber.change1.ex
```
defmodule Jobber do
  alias Jobber.{JobRunner, JobSupervisor}

  def start_job(args) do
    DynamicSupervisor.start_child(JobRunner, {JobSupervisor, args})
  end
end
```

We can now start IEx and repeat our last experiment:

```
iex(1)> Process.whereis(Jobber.JobRunner)
#PID<0.155.0>

iex(2)> Jobber.start_job(work: doomed_job)
{:ok, #PID<0.160.0>}
21:24:34.925 [error] GenServer #PID<0.161.0> terminating
** (RuntimeError) Boom!

-- errors omitted --
```

```
iex(3)> Process.whereis(Jobber.JobRunner)
#PID<0.155.0>
```

The JobRunner process is intact, which means that we successfully isolated potential errors. You also saw how easy it is to start building fault-tolerant applications, by thinking in processes and extending your system when needed.

However, we left one question unanswered. We used a setting called :strategy with a value of :one_for_one when creating the DynamicSupervisor and Supervisor processes. When creating a supervisor from scratch, you have the following supervisor strategies available:

- :one_for_one
- :one_for_all
- :rest_for_one

They tell the supervisor how to manage child process that fail. The one we used so far—:one_for_one—is great when each child process can be restarted independently, without impacting other sibling processes. This is in contrast to the :one_for_all strategy, which will restart *all* child processes when one of them fails. The last one, :rest_for_one, works similarly to :one_for_all, but only sibling processes started *after* the failed one will be restarted. This figure illustrates the differences between the three settings:

:one_for_one

:one_for_all

:rest_for_one

Which strategy to use depends on how your application works. You have to decide how the system should react on error, taking into consideration possible dependencies between processes.

Our job processing system is already looking pretty good, but we are going to add one final feature. Sometimes GenServer processes could be doing intensive work, so it could be useful to limit the number of certain kind of jobs running concurrently. For example, you may want to restrict data imports from a third party system. We already support having a unique id for each Job process, so we just need to learn about process naming to make this work.

Naming Processes Using the Registry

Until now, we used the process identifier we obtained from starting the process to interact with it. However, we can optionally provide a :name value when starting a process, which can be used later on instead of a PID. This way you don't have to keep track of the PIDs of the processes you're starting. Functions like GenServer.cast/2 and GenServer.call/3 also accept a name instead of a process

identifier, and you've already seen how to use Process.whereis/1 to find out the PID for JobRunner.

To make working with names easier, Elixir comes with module called Registry. It is a very efficient key-value store with some useful functionality. Let's see how we can use it.

Processes can be named when they start via the :name option, which is what we did when creating JobRunner:

```
children = [
  {DynamicSupervisor, strategy: :one_for_one, name: Jobber.JobRunner},
]
```

Typically, the :name is set when we link a process using start_link/1. In fact, GenServer.start_link/1 will give each process a name by default, using the name of the module that provides the implementation.

There are three types of accepted values when naming a process:

- an atom, like :job_runner. This includes module names, since they're atoms under the hood, for example Jobber.JobRunner
- a {:global, term} tuple, like {:global, :job_runner}, which registers the process globally. Useful for distributed applications
- a {:via, module, term} tuple, where module is an Elixir module that would take care of the registration process, using the value term

Using Elixir atoms may seem convenient, but could be a problem when starting many uniquely named process, like we do for jobber. Atoms are not garbage collected by the Erlang VM, and there are soft limits on how many atoms could be used in the system.

This is where the Elixir Registry comes in. It allows us to use strings, rather than atoms, which don't have the same limitation, and are generally much easier to work with. The Registry also comes with some helpers for looking up process by name and filtering the results, which are nice additions.

Each Registry works as a process which you have to include in your application supervision tree. You can even start multiple registries if you want to. We are going to create a new Registry next, then use it for registering Job processes dynamically.

Starting a Registry Process

Creating and starting a registry is easy. Make the following change to application.ex:

```
jobber/lib/jobber/application.change2.ex
def start(_type, _args) do
  job_runner_config = [
    strategy: :one_for_one,
    max_seconds: 30_000,
    name: Jobber.JobRunner
  ]

  children = [
    {Registry, keys: :unique, name: Jobber.JobRegistry},
    {DynamicSupervisor, job_runner_config}
  ]

  opts = [strategy: :one_for_one, name: Jobber.Supervisor]
  Supervisor.start_link(children, opts)
end
```

You can see that adding a Registry is similar to adding other processes. Using the :name argument, we set the name of the process to Jobber.JobRegistry. We also used the :keys setting. The keys are the names of the processes we are going to register. Using this setting, we can enforce all keys to be either :unique or allow duplicates with :duplicate. We set the value to :unique.

You can restart IEx, and the JobRegistry process will be available for you to use. Next, let's change our Job process so it uses the Registry when starting up.

Registering New Processes

We are going to make a few changes to job.ex. First, let's add a helper function, via/2, to job.ex:

```
jobber/lib/jobber/job.change1.ex
defp via(key, value) do
  {:via, Registry, {Jobber.JobRegistry, key, value}}
end
```

This will return a tuple that conforms to the required naming specification. The format is always {:via, Registry, config}, where config is also a tuple and can be either {registry_process, key} or {registry_process, key, value}. We will use the extra value element to store some useful metadata about the process. We will label each job with a type description, to make it easier to find out which process is doing what.

You can use the provided id, or generate one, for the name of the process. Let's move some of the existing logic in start_link/1:

```
jobber/lib/jobber/job.change1.ex
def start_link(args) do
  args =
    if Keyword.has_key?(args, :id) do
```

```
      args
    else
      Keyword.put(args, :id, random_job_id())
    end

  id = Keyword.get(args, :id)
  type = Keyword.get(args, :type)

  GenServer.start_link(__MODULE__, args, name: via(id, type))
end
```

We've added an extra :name argument to GenServer.start_link/2 to name the process. This is where we have to provide the :via tuple, generated by the via/2 helper function. We're also passing down the :type argument, which we're going to use later.

We no longer need to generate a new id in the init/1 function, since this is now done in start_link/2. Let's update this:

jobber/lib/jobber/job.change1.ex
```
def init(args) do
  work = Keyword.fetch!(args, :work)
➤  id = Keyword.get(args, :id)
  max_retries = Keyword.get(args, :max_retries, 3)

  state = %Jobber.Job{id: id, work: work, max_retries: max_retries}
  {:ok, state, {:continue, :run}}
end
```

This is all we need to register the process using JobRegistry. Now we can start using the data stored in the registry to put a limit on certain jobs.

Querying the Registry

The Registry module gives you a handy function—select/2. This function uses a matching syntax from Erlang, so it looks a bit cryptic at first. However, once you understand the basic elements of it, you'll have no problem using it.

Let's add a helper function to jobber.ex to retrieve all running processes labeled with import. Then we will break it down and explain what it does:

jobber/lib/jobber.change2.ex
```
def running_imports() do
  match_all = {:"$1", :"$2", :"$3"}
  guards = [{:"==", :"$3", "import"}]
  map_result = [%{id: :"$1", pid: :"$2", type: :"$3"}]
  Registry.select(Jobber.JobRegistry, [{match_all, guards, map_result}])
end
```

Registry.select/2 takes a registry process as its first argument, and a list of match specifications as the last argument. The match specification we have is a bit lengthy, so we have broken it down to three variables:

- match_all is a wildcard that matches all entries in the registry
- guards is a filter that filters results by the third element in the tuple, which has to be equal to "import"
- map_result is transforming the result by creating a list of maps, assigning each element of the tuple to a key, which makes the result a bit more readable

Each value in the Registry is a tuple in the form of {name, pid, value}. In our case, name is the id, and value is the type label. Each element in the tuple is given a special identifier based on the its position in the tuple. So :"$1" corresponds to the first element, which is the id, :"$2" is the pid, and so on. We use these as template variables to match, filter, and map entries in the Registry.

Let's see this in action. Go to .iex.exs and increase the timer for Process.sleep/1 on good_job:

```
good_job = fn ->
  Process.sleep(60_000)
  {:ok, []}
end
```

The function will now take one minute to complete, giving us more time to test things properly. Next, restart IEx for the changes to take effect. We will run three concurrent jobs, and two of them will be import jobs:

```
iex(1)> Jobber.start_job(work: good_job, type: "import")
{:ok, #PID<0.161.0>}

iex(2)> Jobber.start_job(work: good_job, type: "send_email")
{:ok, #PID<0.165.0>}

iex(3)> Jobber.start_job(work: good_job, type: "import")
{:ok, #PID<0.168.0>}
```

Running Jobber.running_imports() should produce something similar to this:

```
iex(4)> Jobber.running_imports()
[
  %{id: "oNnA3I8", pid: #PID<0.161.0>, type: "import"},
  %{id: "OtHGZ6A", pid: #PID<0.168.0>, type: "import"}
]
```

We now have a way to query JobRegistry and limit concurrency for import jobs.

> **Learning about match specifications**
>
> To learn more about match specifications for Registry.select/2, which are also useful when working with ETS, you can check the official Elixir documentation on HexDocs.[3] If you want to dig deeper, Erlang's match specification guide[4] is the best reference on the topic.

Limiting Concurrency of Import Jobs

We are going to put our new running_imports/0 to use. Let's add this check to Jobber.start_job/1:

```
jobber/lib/jobber.change2.ex
def start_job(args) do
  if Enum.count(running_imports()) >= 5 do
    {:error, :import_quota_reached}
  else
    DynamicSupervisor.start_child(JobRunner, {JobSupervisor, args})
  end
end
```

Try starting more than five jobs. After the fifth attempt, you should see this error result:

```
iex(6)> Jobber.start_job(work: good_job, type: "import")
{:error, :import_quota_reached}
```

We have successfully enforced a limit on how many import jobs can run at any given time. We also demonstrated how you can register and lookup processes using the Registry, including adding metadata to each registered process. You can use this in a variety of ways to efficiently manage a large number of processes at runtime.

Our job processing system is now complete. We're going to leave jobber here, but if you'd like to continue working on it, there are many other features you can consider adding, such as job queues or database persistence. You can use this project as a playground to continue practicing building fault-tolerant systems with GenServer and supervisors.

Wrapping Up

Completing this chapter was not an easy challenge, so well done! Learning how to use GenServer correctly will open the doors to building highly-concurrent

---

3. https://hexdocs.pm/elixir/Registry.html#select/2
4. http://erlang.org/doc/apps/erts/match_spec.html

and performant applications, without the complexity you've seen in other programming languages. GenServer is a versatile tool and you will reach for it often when doing work concurrently.

In this chapter, you learned about the most common callbacks used when implementing the GenServer behaviour. You saw how to use each callback, interact with the process while it's still running, and use the process state. You can start processes at runtime using DynamicSupervisor, or build your own supervisor when needed. Finally, you saw how the Registry can make registering and managing even large numbers of running processes a breeze.

The last thing we added to our job processing system was the ability to limit concurrency. This is because running a high volume of concurrent jobs can potentially overwhelm our system, especially when the work requires a lot of system resources. In the next chapter, rather than limiting concurrency, we are going to see how to manage back-pressure and keep workload under control, while maximizing performance at the same time.

Data Processing Pipelines with GenStage

In previous chapters, we covered several approaches to executing data asynchronously. Although different in their own way, they had one thing in common: we were deciding the amount of work that had to be done and Elixir would then eagerly process the work to give us the result.

This has some potential drawbacks. For example, we have a finite amount of memory and CPU power available. This means that our server may become overwhelmed by the amount of work it needs to do and become slow or unresponsive. Often we rely on third-party API services, which have rate limiting in place and fixed quotas for the number of requests we can make. If we go over their quota, requests will be blocked and our application will stop working as expected.

In a nutshell, as the amount of work we need to do grows, there is an increasing chance that we can hit a certain limit on a resource available to us. You can spend more money on computer hardware or buy more powerful virtual machines, but that's usually a temporary solution that often leads to diminishing returns. For that reason, making the best use of existing resources should be a priority when designing and scaling software applications.

In this chapter, we're going to learn how to build data processing pipelines that can utilize our system resources reliably and effectively. We'll introduce the GenStage Elixir library and the fundamental building blocks that it provides. First we are going to create a simple data processing pipeline to start with, then scale it and extend it to demonstrate how you can tackle more complex use cases.

But before we get into it, first we need to explain what *back-pressure* is, and how it enables us to build data processing pipelines. Let's get started!

Understanding Back-Pressure

Imagine you're a famous writer giving autographs at a book event. There is a crowd of people rushing to meet you. You can only do one autograph at a time, so the organizers let people in one by one. When you sign someone's book, you ask for the next person to come forward. What if the organizers suddenly let everyone in? Of course, it will be complete chaos! You'll try to sign everyone's books as quickly as you can but soon you'll get stressed and exhausted, leaving the event early.

It is much more efficient to have an orderly queue of people and to take your time to sign each book. Because you always ask for the next person to come forward, you are in control of the amount of work you have to do and it is much easier to keep going. Maybe you won't get that tired, so you decide to stay at the event longer and make sure everyone gets an autograph. This is in fact an example of handling back-pressure in real life.

How does this translate in programming? Using the GenStage library we are going to build a data processing pipeline that works like the well-organized book event we just described. The system will process only the amount of work it can handle at a given time, just like the famous writer from our example. If the system has free capacity it will politely ask for more work and wait for it.

This simple shift in our thinking is very powerful. It enables us to build complex data pipelines that regulate themselves to utilize the available resources in the best possible way.

Borrowing terminology

The term *back-pressure*, according to Wikipedia, originates from fluid dynamics and the automotive industry, where it is used to describe resistance to the normal flow of fluids in pipes.

Software engineers borrowed the term and loosely use it in the context of data processing, when something is slowing down or stopping the flow of data. When we talk about using back-pressure, we actually mean using a mechanism to control or handle back-pressure in a beneficial way.


Introducing GenStage

GenStage was originally developed by José Valim, the creator of Elixir, and released in July 2016. As he described it in the official announcement:

"GenStage is a new Elixir behaviour for exchanging events with back-pressure between Elixir processes."

In the previous chapter, we used the GenServer behaviour to build long-running *server* processes. The GenStage behaviour, as its name suggests, is used to build *stages*. Stages are also Elixir processes and they're our building blocks for creating data processing pipelines.

Stages are simple but very powerful. They can receive *events* and use them to do some useful work. They can also send events to the next stage in the pipeline. You can do that by connecting stages to each other, creating something like a chain, as you can see here:


A stage can also have multiple instances of itself. Since stages are processes, this effectively means running more than one process of the same stage type. This means that you can also create something that looks like this:


These are just two possible designs. In practice, different problems require different solutions, so your data processing pipeline could end up looking completely different. When you finish this chapter you will have a solid understanding of GenStage and you will be able to come up with a solution that works best for you.

As you can see, stages are very flexible and can be used in a variety of ways. However, their most important feature is back-pressure.

Although events move between stages from left to right on our diagram, it is actually the last stage in the pipeline that controls the flow. This is because the *demand* for more events travels in the opposite direction—from right to left. This figure shows how it works:

Stage D has to request events from Stage C, and so on, until the demand reaches the beginning of the pipeline. As a rule, a stage will send demand for events only when it has the capacity to receive more. When a stage gets too busy, demand will be delayed until the stage is free, slowing down the flow of data. Sounds familiar? This is exactly how the book event example from Understanding Back-Pressure, on page 58 would work if we implemented it as a data processing pipeline. This figure illustrates how we can model the book event using stages:


By treating guests as a flow of events, the author stage can "process" a guest by signing their book, and request the next one from the organizer when

ready. The organizer stage, on the other hand, has to make sure that only one guest at a time comes forward, as the author requests. This is how back-pressure works in GenStage, and the best part is that you benefit from it simply by thinking in stages and connecting them together.

Connecting stages is easy, but first you need to know what stages to use. There are three different types of stages available to us: producer, consumer, and producer-consumer. Each one plays a certain role. Let's briefly cover each type of stage and see how it is made to work with the rest.

The Producer

At the beginning of a data processing pipeline there is always a producer stage, since the producer is the source of data that flows into the pipeline. It is responsible for producing *events* for all other stages that follow. An *event* is simply something that you wish to use later on, it could be a map or a struct. You can use any valid Elixir data type. Here is an example of a two-stage data pipeline:


The events produced by the producer have to be processed by another stage, which is missing in the figure. What we need here is a consumer for those events.

The Consumer

Events created by the producer are received by the consumer stage. A consumer has to *subscribe* to a producer to let them know they're available, and request events.

This producer and consumer relationship is around us every day. For example, when you go to the farmer's market, the farmers that grow vegetables are producers. They wait until a customer (a consumer) comes and asks to buy some vegetables.

Consumer stages are always found at the end of the pipeline. Now we can fill in the missing process. The following figure illustrates the simplest data pipeline we can create using GenStage:

It has just a single producer and a consumer.

The Producer-Consumer

Although having a producer and a consumer is already very useful, sometimes we need to have more than two stages in our pipeline. This is where the producer-consumer stage comes in—it has the special ability to produce and consume events at the same time. The following figure shows how one or more producer-consumer stages can be used to extend a pipeline.

A useful analogy to a producer-consumer is the restaurant. A restaurant serves meals (producer) but in order to cook the meals, it needs ingredients sourced from its suppliers (acting as a consumer). In a nutshell, producer-consumers are the middle-man in our data processing pipelines.

We covered a lot of theory, so let's finally dive into some examples.

Building Your Data Processing Pipeline

Complex use cases may require a data processing pipeline with a consumer stage, one or more producers, and several producer-consumers in-between. However, the main principles stay the same, so we're going to start with a 2-stage pipeline first and demonstrate how that works.

We will build a fake service that scrapes data from web pages—normally an intensive task, dependent on system resources and a reliable network connection. Our goal is to be able to request a number of URLs to be scraped, and have the data pipeline take care of the workload.

First, let's create a new application with a supervision tree, like we've done before. We will name it scraper and pretend we're going to scrape data from web pages.

```
$ mix new scraper --sup
```

This will print the usual information in the console. Let's navigate in the new directory and edit mix.exs to add gen_stage as dependency:

scraper/mix.exs
```
defp deps do
  [
    {:gen_stage, "~> 1.0"}
  ]
end
```

Next, run the command mix do deps.get, compile to download and compile all dependencies.

Let's add a dummy function to scraper.ex which will simulate doing some time-consuming work:

scraper/lib/scraper.ex
```
def work() do
  # For simplicity, this function is
  # just a placeholder and does not contain
  # real scraping logic.
  1..5
  |> Enum.random()
  |> :timer.seconds()
  |> Process.sleep()
end
```

The work/0 function will take a random number from 1 to 5, convert it from seconds to milliseconds using Erlang's :timer.seconds/1 and then pass it to Process.sleep/1, pausing the process for that amount of time. As a result, calling this function will pause the current process from 1 to 5 seconds at random.

Now, let's create our producer.

Creating a Producer

Our scraper will scrape data from web pages. At the beginning of our pipeline we will put a producer stage that will be responsible for what web pages we are going to scrape. In other words, it will *produce* URLs for consumers interested in getting web pages to scrape.

We are going to name our producer PageProducer and define it in a new file page_producer.ex in the lib directory:

```
scraper/lib/page_producer.ex
defmodule PageProducer do
  use GenStage
  require Logger

  def start_link(_args) do
    initial_state = []
    GenStage.start_link(__MODULE__, initial_state, name: __MODULE__)
  end

  def init(initial_state) do
    Logger.info("PageProducer init")
    {:producer, initial_state}
  end

  def handle_demand(demand, state) do
    Logger.info("Received demand for #{demand} pages")
    events = []
    {:noreply, events, state}
  end
end
```

The code above looks similar to a GenServer process code. Instead of use GenServer there is use GenStage. There is a start_link/1 function and an init/1 function. Of course there are some differences, too.

The init/1 function for GenStage processes must specify what type of stage the process is. Since we want to create a producer, we return {:producer, initial_state}. The first element in the tuple—the stage type—could be :producer, :consumer or :producer_consumer; the second is the argument that sets the initial state for the process.

We are also implementing a handle_demand/2 callback, which you haven't seen before. This callback is required for processes of type :producer and :producer_con-sumer. When a :consumer process asks for events, handle_demand/2 will be invoked with two parameters: the number of events requested by the consumers and the internal state of the producer. In the result tuple, the second element must be a list containing the actual events. Right now we are just returning an empty list, but we will re-visit this part later.

Creating a Consumer

Now that there is a producer in place, let's create a consumer. Since we name our producer PageProducer, we are going to call the consumer PageConsumer. Create the file page_consumer.ex and place it in the same lib directory.

```
scraper/lib/page_consumer.ex
defmodule PageConsumer do
  use GenStage
```

```elixir
require Logger

def start_link(_args) do
  initial_state = []
  GenStage.start_link(__MODULE__, initial_state)
end

def init(initial_state) do
  Logger.info("PageConsumer init")
  {:consumer, initial_state, subscribe_to: [PageProducer]}
end

def handle_events(events, _from, state) do
  Logger.info("PageConsumer received #{inspect(events)}")

  # Pretending that we're scraping web pages.
  Enum.each(events, fn _page ->
    Scraper.work()
  end)

  {:noreply, [], state}
end
end
```

The code looks familiar and somewhat similar to a producer. What has changed is the atom returned by init/1, specifying the stage type. This time we are returning :consumer.

We also return some options from the init/1 function. We are telling this consumer to *subscribe* to PageProducer, via subscribe_to: [PageProducer]. This is how stages are linked together.

Since consumers receive events, we have to implement a special callback for stages of type :consumer and :producer_consumer, and this is the handle_events/3 callback. When a consumer asks for events, and the producer responds, the list of events given to a consumer will be the first parameter in the handle_events/3 function. This is where we would do something with the events and use it in a meaningful way. For now, we're just going to call our Scraper.work/1 dummy function.

Subscribing at runtime

You can programmatically subscribe to a producer at runtime using sync_subscribe/3 and async_subscribe/3 from the GenStage module. This is useful when consumers and/or producers are created dynamically at runtime. Note that you will also have to handle the re-subscribe yourself if the producer crashes and is restarted.

Now that we have a producer and a consumer, let's add them to our application supervision tree in application.ex. Our start/2 function should look like so:

```
scraper/lib/scraper/application.ex
def start(_type, _args) do
  children = [
    PageProducer,
    PageConsumer
  ]

  opts = [strategy: :one_for_one, name: Scraper.Supervisor]
  Supervisor.start_link(children, opts)
end
```

Here we are introducing another type of child specification. This format is the shortest, and is equivalent to this:

```
children = [
  {PageProducer, []},
  {PageConsumer, []}
]
```

We are going to use this format when we don't need to pass any options to the process we are adding.

Running the app with iex -S mix should produce this:

```
19:28:18.351 [info]  PageProducer init
19:28:18.359 [info]  PageConsumer init
19:28:18.359 [info]  PageProducer received demand for 1000 pages
```

It seems that the producer and consumer are talking to each other—success. However, there are two issues: first, the producer doesn't actually produce anything (yet); and second, the consumer asked for a whopping 1000 pages. Let's tackle the issues in reverse order.

Understanding Consumer Demand

You saw that PageProducer was asked for 1000 events and you may wonder where the number comes from. By default, stages of type :consumer and :producer_consumer make sure that demand for new events is between 500 and 1000. You can configure this through the min_demand and max_demand settings. When a consumer subscribes to a producer in the init/1 function, you can pass a tuple with some configuration options:

```
# `subscribe_to` options can be a list of tuples.
sub_opts = [{PageProducer, min_demand: 500, max_demand: 1000]
{:consumer, initial_state, subscribe_to: sub_opts}
```

By configuring demand you can control two things: max_demand will set the maximum number of events the consumer can ask for, and min_demand will

determine the lowest number of events available to the consumer before it can ask for more.

The consumer is greedy and will immediately demand 1000 events on start, according to the max_demand value you set. Let's say the producer supplies 1000 events. The consumer will process the first batch using a simple formula:

events to process = max_demand - min_demand

This is to ensure that while the consumer is processing the events, fresh demand is issued to the producer to keep it busy supplying the next batch of events.

Since max_demand is 1000 and min_demand is 500 by default, the consumer will process 500 events first, and then the remaining 500, according to the formula. This means that the handle_events/3 callback will be called with a list of 500 events initially, followed by another 500 when it is done processing the previous one. The final outstanding demand for the producer will be 500 events, as shown in this figure:
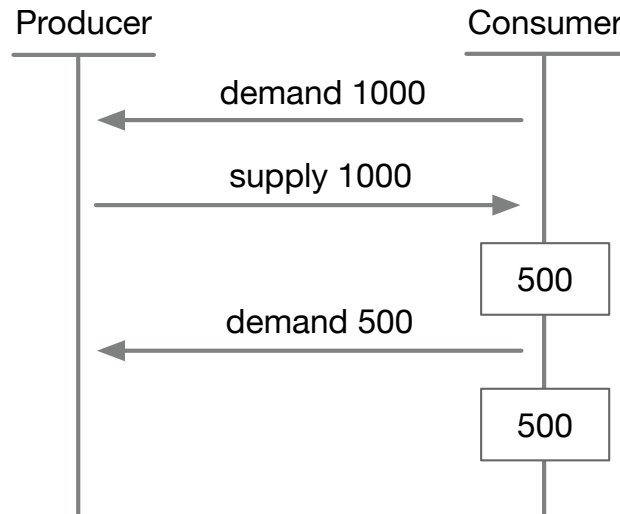
Let's take another example with the same settings. This time the producer responds with 600 events, even though the consumer asked for 1000. According to the formula, the 600 events will be broken down to batches of 500 and 100. While the consumer is processing the first batch of 500 events, there is still an outstanding demand for 400 events from the initial request. As soon as the consumer starts processing the first batch, it will ask for 500

events more. This will bring total consumer demand to 900 events. This figure illustrates the example:
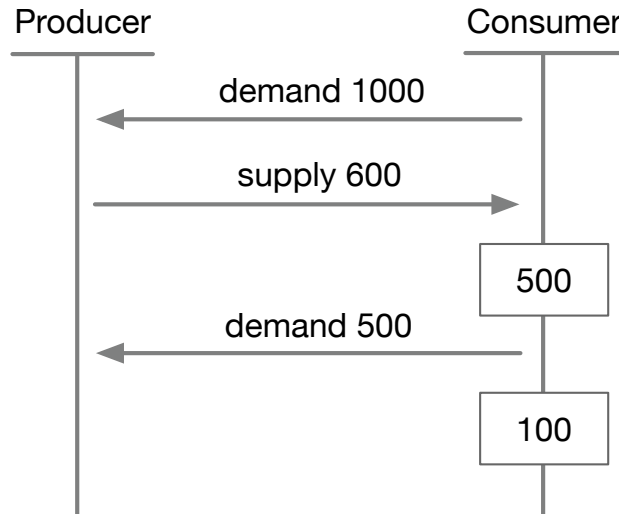
This sounds complicated and confusing at first, especially since we are dealing with concurrent processes, but there are good reasons behind this implementation. By tracking the number of events a consumer processes and number of events in progress, GenStage optimizes the flow of data for peak efficiency, keeping both consumers and producers busy at all time. The consumer demands more events only when a batch has been processed, so if the processing of the events takes longer, demand will be delayed. The opposite is also true—if events are quick to process, demand will be issued sooner. In all cases, demand will be kept within the given min_demand and max_demand values and will be constantly adjusted to fit this range by GenStage.

You don't have to always process events in batches. If you lower the value of max_demand to 1, the events will be requested one by one. Depending on your use case, you have to come up with the configuration that works best for you. The default values of 500 and 1000 are simply there to get you started. This could be challenging, so here's a simple questionnaire that will give you a good starting point:

• How long does it take to process a single event vs 100 vs 1000?

  – This will help you determine if it is faster to demand a batch of events, and what batch size (as opposing to one by one);
• How often do you expect new events?

- If the producer produces events frequently and they're in constant supply, it may be beneficial to increase min_demand, especially when working in batches is ideal;
- Can processing be delayed?
  - Since consumers will wait until min_demand is satisfied by the producer, you will want to lower min_demand or even set it to 0 (demanding an event as soon as it is available by the producer);
- What system resources are available to me?
  - You can distribute the load by lowering min_demand and max_demand, buffering events or using several consumers (more about these later);

Answering these questions should help you get started with your own GenStage implementation. It may also take you more than one attempt to find the best min_demand and max_demand strategy, so don't despair. It is always a good idea to measure how your pipeline performs and tweak the demand until you're happy with the results.

Now, for our scraper project, let's lower the demand to something more suitable:

```
scraper/lib/page_consumer.change1.ex
def init(initial_state) do
  Logger.info("PageConsumer init")
  sub_opts = [{PageProducer, min_demand: 0, max_demand: 3}]
  {:consumer, initial_state, subscribe_to: sub_opts}
end
```

With min_demand set to 0 and max_demand set to 3, we are ensuring that the consumer will take at least 1 event when it is available, up to 3 at a time. However, our producer still has to supply the consumers with events, and right now it doesn't do anything. Let's fix this.

Re-Visiting the Producer

Going back to PageProducer, we have implemented the handle_demand/2 callback to respond to demand and produce events. Remember that handle_demand/2 is always called when the producer receives demand from a consumer. Since consumers always demand events when they're free, this callback will be useful when we want to respond to consumer demand immediately and keep consumers busy.

Instead, we would like to expose an API which developers can use to request pages to be scraped. In this particular use case, consumers will have to wait patiently until work is available for them. Therefore, handle_demand/2 won't be very useful for what we're trying to achieve.

Thankfully, there are several ways for a producer to dispatch events. In fact, most callbacks for :producer and :producer_consumer stages have the ability to do so.

Earlier we mentioned that GenStage is built on top of GenServer. This means that the callbacks we covered in Callbacks In-Depth, on page 27 are also available for GenStage processes:

- handle_call/3
- handle_cast/2
- handle_info/3

They will be called when you invoke GenStage.call/3, GenStage.cast/2 or Process.send/3, respectively. However, the return signatures of those callbacks have an important difference to their GenServer counterparts. Here are two examples of return tuples allowed for GenStage:

```
{:reply, reply, [event], new_state}
{:noreply, [event], new_state}
```

Notice the extra element in the tuple that holds a list of event values. These callbacks work exactly the same as the handle_demand/3 callback. This is great news, because it gives us a lot of flexibility when dispatching events.

Let's implement our API in PageProducer:

scraper/lib/page_producer.change1.ex
```
def scrape_pages(pages) when is_list(pages) do
  GenStage.cast(__MODULE__, {:pages, pages})
end

def handle_cast({:pages, pages}, state) do
  {:noreply, pages, state}
end
```

We have exposed a function scrape_pages/1 which accepts a list of URLs. This function will be our user-facing API. Inside the function, we call GenStage.cast/2, just like we did with GenServer before.

In the handle_cast/2 callback function, we return a tuple as a result. The first element in the tuple is always the type of reply, which is :noreply in this case. The second element must be a list, containing the events we want to dispatch. We are going to return pages, which contains our list of strings. Finally, the third element is the process state, which you can update if you need to. Here, we are just returning it unchanged. As you can see, the return tuple format is very similar to the one for GenServer, but with the addition of the events element.

Let's run our application again using the IEx shell:

```
$ iex -S mix
```

You should see an output similar to this:

```
Erlang/OTP 21 [erts-10.0.3] [source] [64-bit] [smp:4:4] [ds:4:4:10]
  [async-threads:1] [hipe] [dtrace]

Compiling 1 file (.ex)

16:10:30.437 [info]  PageProducer init
16:10:30.443 [info]  PageConsumer init
16:10:30.443 [info]  PageProducer received demand for 3 pages

Interactive Elixir (1.8.1) - press Ctrl+C to exit (type h() ENTER for help)

iex(1)>
```

As expected, the consumer sends demand as soon as it is initialized. Since our handle_demand/2 callback does not return events, this initial demand is not satisfied and therefore the consumer will wait until events are available.

Now, we're going to create a list of URLs and call our API:

```
iex(1)> pages = [
...(1)> "google.com",
...(1)> "facebook.com",
...(1)> "apple.com",
...(1)> "netflix.com",
...(1)> "amazon.com"
...(1)> ]
["google.com", "facebook.com", "apple.com", "netflix.com", "amazon.com"]

iex(2)> PageProducer.scrape_pages(pages)
```

Let's look closely at the output log:

```
16:19:51.733 [info]  PageConsumer received ["google.com",
  "facebook.com", "apple.com"]
16:20:02.742 [info]  PageProducer received demand for 1 pages
16:20:02.743 [info]  PageConsumer received ["netflix.com", "amazon.com"]
```

We can see that PageConsumer immediately received the first three pages, which took a bit of time to process, judging by the timestamps. Since only two pages were available next, our consumer realized that it has capacity for one more page, so it immediately issued demand for another page, while starting work on the other two. That's great, everything is working as expected.

Congratulations, you just created your first data processing pipeline with GenStage! We have created a producer and a consumer, put them to work together and introduced a lot of new concepts on the way.

However, we are not finished with our scraper project just yet. Although it is working correctly, we can in fact easily improve it with just a few small changes. Since our Scrape.work/0 function is very slow, working in batches and scraping pages in a sequence is not ideal. We can significantly improve performance by taking advantage of modern multi-core CPUs to scrape pages concurrently. Let's see how that works.

Adding More Consumers

Rather than working in batches within a single process, we can easily scale our data processing pipeline by running more than one consumer, each responsible for scraping one page at a time. To do this, let's adjust PageConsumer's demand:

```
scraper/lib/page_consumer.change2.ex
def init(initial_state) do
  Logger.info("PageConsumer init")
➤  sub_opts = [{PageProducer, min_demand: 0, max_demand: 1}]
  {:consumer, initial_state, subscribe_to: sub_opts}
end
```

and add another PageConsumer to our supervision tree:

```
scraper/lib/scraper/application.change1.ex
children = [
  PageProducer,
➤  Supervisor.child_spec(PageConsumer, id: :consumer_a),
➤  Supervisor.child_spec(PageConsumer, id: :consumer_b)
]
```

Now our consumer will take only one event at a time, but we have two consumer processes running concurrently. As soon as one is free, it will issue demand to scrape another page.

Notice that when we added another PageConsumer, we used Supervisor.child_spec/2. As we saw, each process should have a unique ID in the supervision tree. In the example above, the processes are called :consumer_a and :consumer_b. If we do not do that, we will get an error as soon as the main supervisor initializes. We can also use the Registry module to assign a name to each process, as we have done in Naming Processes Using the Registry, on page 49.

With this approach, we can add as many consumer processes as needed and GenServer will distribute the events for us, acting as a load balancer. Let's try the scrape_pages/1 function again and compare the results:

```
16:44:58.617 [info]  PageConsumer received ["facebook.com"]
16:44:58.617 [info]  PageConsumer received ["google.com"]
16:45:01.618 [info]  PageConsumer received ["apple.com"]
```

```
16:45:02.618 [info]  PageConsumer received ["netflix.com"]
16:45:02.619 [info]  PageConsumer received ["amazon.com"]
16:45:05.620 [info]  PageProducer received demand for 1 pages
16:45:06.619 [info]  PageProducer received demand for 1 pages
```

If you look closely at the timestamps, you can see how the first two pages were immediately picked up by the waiting consumers; since the work takes a variable amount of time, the first consumer to finish takes the next URL off the list of events. This repeats until all work is done and the consumer issues demand for events once again.

This is a great improvement. Not only we have processed the pages faster, but we now have confidence that we can scale our pipeline when we need to. You can call scrape_pages/1 even when the consumers are still busy, so events will be queued up automatically. It is important to understand how this works, so we're going to briefly cover this next.

Buffering Events

At this point, you may suspect that producers keep dispatched events in memory—and you will be correct. Producers have a built-in buffer which is used whenever the number of dispatched events is greater than the total pending demand. As we saw earlier, events staying in the buffer are automatically granted to consumers who issue demand.

The default size of the buffer is 10,000 events for stages of type :producer, and :infinity for type :producer_consumer. However, both can be configured with a fixed capacity of our choice or :infinity. In the init/1 callback, we can provide the optional buffer_size parameter in the return tuple:

```
{:producer, initial_state, buffer_size: :infinity}
```

Now, let's do a quick experiment. Change buffer_size to 1 and re-run the application. Call scrape_pages/1 with the same list of pages as before. The output this time will be different.

```
iex(1)> pages = ["google.com", "facebook.com",
  "apple.com", "netflix.com", "amazon.com"]

iex(2)> PageProducer.scrape_pages(pages)

[info]  PageConsumer received ["facebook.com"]
[info]  PageConsumer received ["google.com"]
[warn]  GenStage producer PageProducer has
  discarded 2 events from buffer
[info]  PageConsumer received ["amazon.com"]
[info]  PageProducer received demand for 1 pages
[info]  PageProducer received demand for 1 pages
```

We did this to demonstrate what happens when the buffer_size limit is exceeded. We received a warning, telling us how many events were dropped from the buffer and for which producer. Since demand by the two consumers was two events (one each), and we dispatched five events, we ended up with three extra events and no demand for them. Because buffer_size is 1, the first two events were dropped from the list. Only the last one was kept in the buffer and processed later.

Dropping events from the end of the buffer

If you want to use a fixed-size buffer, you also have the option to discard events from the end of the queue when the :buffer_size limit is hit. Just pass the optional :buffer_keep param and set it to :first (the default value is :last).

Using the built-in buffer is convenient for most use cases. If you need fine-grain control over the number of events produced and dispatched, you may want to look into implementing your own queue for storing produced events and pending demand. Erlang's :queue[1] is a great option as it is already available in Elixir. Such a queue could be stored in producer's state, and used to dispatch events only when demand has been registered in handle_demand/3. This will also give you an opportunity to implement your custom logic for discarding events—useful if you want to prioritize one type of event over another.

Adding Concurrency With ConsumerSupervisor

Using multiple consumers to process events one by one, concurrently, is a very useful technique. To make this task even easier, GenStage comes with a special type of supervisor called ConsumerSupervisor. It works similarly to a consumer and can subscribe to one or more producers. It can also monitor, start and re-start child processes, just like a supervisor.

What makes ConsumerSupervisor special is that when it receives a list of events from a producer, it automatically starts a process for each event and passes the event as an argument to that process. When a child process exits successfully, new demand will be issued by ConsumerSupervisor and the cycle repeats. This figure illustrates how child processes are started on-demand when new events are received:

---

1. https://erlang.org/doc/man/queue.html
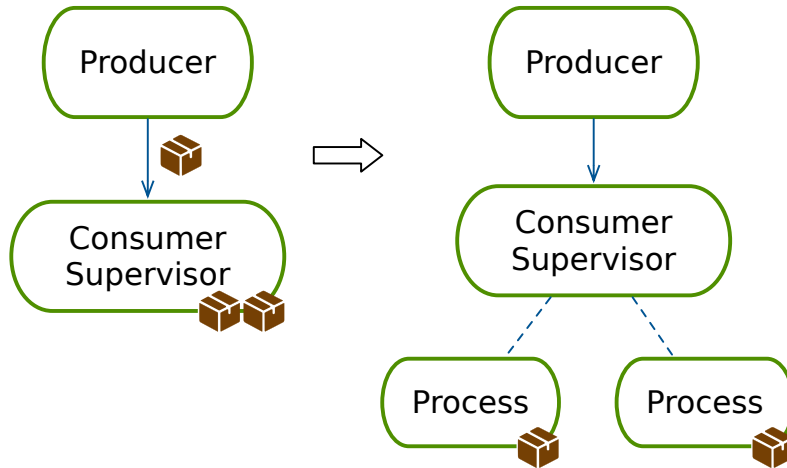
Next, we are going to refactor some of our existing logic using a ConsumerSupervisor to demonstrate how it works in practice.

Creating a ConsumerSupervisor

We're going to create a new file page_consumer_supervisor.ex and place it in the lib directory. Here are the full contents:

```
scraper/lib/page_consumer_supervisor.ex
defmodule PageConsumerSupervisor do
  use ConsumerSupervisor
  require Logger

  def start_link(_args) do
    ConsumerSupervisor.start_link(__MODULE__, :ok)
  end

  def init(:ok) do
    Logger.info("PageConsumerSupervisor init")

    children = [
      %{
        id: PageConsumer,
        start: {PageConsumer, :start_link, []},
        restart: :transient
      }
    ]

    opts = [
      strategy: :one_for_one,
      subscribe_to: [
        {PageProducer, max_demand: 2}
      ]
```

```
    ]
    ConsumerSupervisor.init(children, opts)
  end
end
```

This is a lot of new code, so let's break it down and explain how it works.

We named our module PageConsumerSupervisor and immediately after the defmodule declaration, we brought in the ConsumerSupervisor module logic. Since this is a process, we define the now familiar start_link/1 function. State is not relevant for ConsumerSupervisor, so we simply pass an :ok atom.

The init/1 function is more interesting. Here, we have to delegate the initialization to ConsumerSupervisor.init/2, which accepts two parameters, in this order:

- a list of child specifications. ConsumerSupervisor needs only one child spec, with a mandatory :restart key. The only restart options supported are :temporary and :transient;
- a list of arguments. Here you have to configure the subscription, using the :subscribe_to key, which works like the same option in the stages of type :consumer and :producer_consumer. We also set a supervisor strategy using the :strategy key, just like with a standard Supervisor process.

For first parameter—the list of child specifications—we have created a children variable, which contains the specification to start the PageConsumer module. As we said already, we only need one child spec here. This is because Consumer-Supervisor will automatically replicate this child spec when starting new processes. You can think of it as a template, which will be used for all children processes.

We also have an opts variable holding the supervisor options. We picked the :one_for_one strategy, and we have also subscribed to PageProducer. Since PageConsumerSupervisor starts child processes for each event, min_demand and max_demand will tell the supervisor how many processes it needs to start, according to the demand. We are going to set max_demand to 2, which means that two consumers (at most) could run concurrently.

Hopefully these options are already familiar to you. If you have forgotten what they mean or have skipped some parts of the book, please refer to Chapter 2, Long-Running Processes Using GenServer, on page 25 for more information.

It may look like we have added a lot of new code, but with this setup, our consumer will become much simpler. Let's move on to that next.

The Simplified Consumer

As we mentioned before, PageConsumerSupervisor will wait for a child process to exit successfully, before it issues new demand. We know from Chapter 1, Easy Concurrency With the Task Module, on page 1 that the easiest way to start a new background process is to use the Task module. The Task.start_link/1 function can run our code in a process and then exit—exactly what we need. Here is our refactored PageConsumer:

```
scraper/lib/page_consumer.change3.ex
defmodule PageConsumer do
  require Logger

  def start_link(event) do
    Logger.info("PageConsumer received #{event}")

    Task.start_link(fn ->
      Scraper.work()
    end)
  end
end
```

PageConsumer is now just a few lines of code. This may be surprising and you probably wonder what happened with all other logic.

Remember that we told PageConsumerSupervisor to subscribe to PageProducer for events. This means that PageConsumerSupervisor has effectively taken the place of a :consumer in our data processing pipeline. However, PageConsumerSupervisor only manages demand, receives events and starts new processes. It doesn't do any work.

That's why we still need PageConsumer to process the events. Our event processing logic is still the same, but it has been moved from handle_events/3 to Task.start_link/1. Notice that we also no longer receive a list of one or more events, just a single event variable.

Although PageConsumer is no longer a GenStage of type :consumer, it is still a consumer of events, so we are going to stick to this convention and keep its name as it is.

Putting It All Together

Finally, we need to make a small change in our main application supervision tree in application.ex. We are going to remove the individual PageConsumer processes and add PageConsumerSupervisor in their place:

```
scraper/lib/scraper/application.change2.ex
children = [
  PageProducer,
```

➤   PageConsumerSupervisor
]

Let's re-run our application and call PageProducer.scrape_pages/1 in the IEx shell, using the same list of URLs:

```
iex(2)> PageProducer.scrape_pages(pages)

14:28:18.783 [info]  PageConsumer received google.com
14:28:18.785 [info]  PageConsumer received facebook.com
14:28:21.790 [info]  PageConsumer received apple.com
14:28:23.790 [info]  PageConsumer received netflix.com
14:28:25.791 [info]  PageConsumer received amazon.com
14:28:26.791 [info]  PageProducer received demand for 1 pages
14:28:27.792 [info]  PageProducer received demand for 1 pages
```

Great, everything works just like before, but we have significantly simplified our code. The supervisor will monitor processes and will restart them in case of an error. Furthermore, it is now even easier to scale our data processing pipeline. We just have to go to PageConsumerSupervisor and increase max_demand.

How to find out number of logical cores at runtime?

In Elixir, you can get the number of logical cores available programmatically by calling System.schedulers_online(). This could be useful if you want to set ConsumerSupervisor's max_demand dynamically, for example:

max_demand = System.schedulers_online() * 2

My personal laptop has a CPU with 4 logical cores, so max_demand will be 8 using the formula above. However, when I deploy the application to a production server, it will run on a much more powerful machine with 8 logical cores. As a result, the maximum number of concurrent processes will be 16—much better. For best results, always measure and adjust max_demand according to your use case.

It's worth highlighting that ConsumerSupervisor can start any kind of process. We picked Task.start_link/1 for simplicity and versatility. However, you can start GenServer processes or even GenStage producer processes if you need to. Keep in mind that all child processes must exit with reason :normal or :shutdown, so the supervisor can re-issue demand. You can do this by returning {:stop, :normal, state} from a process callback when you're ready to terminate it. Most GenServer and GenStage callbacks support this return tuple, including handle_demand/3 and handle_events/3, which gives you a lot of flexibility.

Creating Multi-Stage Data Pipelines

We already demonstrated how :producer and :consumer stages work in practice. The only type of stage that we haven't seen in action yet is the :producer_consumer. Producer-consumer stages are the key to building infinitely complex data processing pipelines. The good news is that if you understand how producers and consumers work, you already know producer-consumers.

When you learn how to add stages and extend your data pipelines, you may be tempted to start organizing your business logic using stages, rather than plain Elixir modules and functions. As the GenStage documentation warns us, this is an anti-pattern:

> If your domain has to process the data in multiple steps, you should write that logic in separate modules and not directly in a GenStage. You only add stages according to the runtime needs, typically when you need to provide back-pressure or leverage concurrency.

A good rule of thumb is to always start with plain functions. When you recognize the need for using back-pressure, create a 2-stage data pipeline first. As we are going to see in a moment, adding more stages is easy, so gradually you can extend it when you spot an opportunity for improvement.

Right, let's begin. First, we need to add some business logic that justifies adding another stage. Open scraper.ex and add the following function:

```
scraper/lib/scraper.change1.ex
def online?(_url) do
  # Pretend we are checking if the
  # service is online or not.
  work()

  # Select result randomly.
  Enum.random([false, true, true])
end
```

We are going to pretend that we are sending a network request to check if the website is online. Our work/1 function comes in handy, so we are going to use it to pause the current process. As the function name suggests, we return true or false, selected at random from three options: [false, true, true]. This way we have 33% chance of the service being offline.

Now that our "business logic" is in place, let's add a new stage to our pipeline.

Adding a Producer-Consumer

We're going to name our new module OnlinePageProducerConsumer, so let's create online_page_producer_consumer.ex in the lib directory.

First, we're going to define the start_link/1 function:

```
scraper/lib/online_page_producer_consumer.ex
def start_link(_args) do
  initial_state = []
  GenStage.start_link(__MODULE__, initial_state, name: __MODULE__)
end
```

The start_link/1 function will take no arguments and have no state. It's important to set the :name key of the process so we can refer to it later. Next, add the init/1 function:

```
scraper/lib/online_page_producer_consumer.ex
def init(initial_state) do
  Logger.info("OnlinePageProducerConsumer init")

  subscription = [
    {PageProducer, min_demand: 0, max_demand: 1}
  ]

  {:producer_consumer, initial_state, subscribe_to: subscription}
end
```

The return tuple of our producer-consumer is very similar to the one we have in PageConsumer. However, notice that the type of stage is now :producer_consumer instead of :consumer.

Finally, we will implement the handle_events/3 callback:

```
scraper/lib/online_page_producer_consumer.ex
def handle_events(events, _from, state) do
  Logger.info("OnlinePageProducerConsumer received #{inspect(events)}")
  events = Enum.filter(events, &Scraper.online?/1)
  {:noreply, events, state}
end
```

To make things a bit more interesting, we are going to filter offline services and only dispatch URLs for websites that are online. That's correct—unlike :consumer stages, we can return a list of events from the handle_events/3 callback.

Producer-consumer stages can also implement handle_demand/3, just like any producer, but we will leave it out for now. It will dispatch no events by default.

The most difficult part is now over. All have to do is plug the new stage into the pipeline.

Re-Wiring Our Pipeline

We want to scrape pages which are online, so PageConsumerSupervisor will no longer subscribe to PageProducer, but to OnlinePageProducerConsumer. Perhaps Page-ConsumerSupervisor and PageConsumer should be renamed to OnlinePageConsumerSuper-

visor and OnlinePageConsumer, respectively, but we will leave them as they are for now.

Let's swap the following part in page_consumer_supervisor.ex:

scraper/lib/page_consumer_supervisor.ex
```
subscribe_to: [
  {PageProducer, max_demand: 2}
]
```

with

scraper/lib/page_consumer_supervisor.change1.ex
```
subscribe_to: [
  {OnlinePageProducerConsumer, max_demand: 2}
]
```

And edit application.ex to add the producer-consumer into the main supervision tree:

scraper/lib/scraper/application.change3.ex
```
children = [
  PageProducer,
➤  OnlinePageProducerConsumer,
  PageConsumerSupervisor
]
```

We added OnlinePageProducerConsumer before PageConsumerSupervisor on purpose. Producers always have to be started before the consumers. If we don't do that, the consumers won't be able to subscribe to them because the producer process would not be available.

That's it. We didn't have to change our external API, so you can start IEx and call PageProducer.scrape_pages/1 with our test list of pages. You should see something like this:

```
12:49:34.111 [info]  OnlinePageProducerConsumer received ["google.com"]
12:49:35.112 [info]  PageConsumer received google.com
12:49:35.112 [info]  OnlinePageProducerConsumer received ["facebook.com"]
12:49:36.113 [info]  PageConsumer received facebook.com
12:49:37.113 [info]  OnlinePageProducerConsumer received ["apple.com"]
12:49:38.114 [info]  OnlinePageProducerConsumer received ["netflix.com"]
12:49:41.115 [info]  OnlinePageProducerConsumer received ["amazon.com"]
12:49:45.116 [info]  PageProducer received demand for 1 pages
```

We have randomized our logic so the output will be slightly different every time. Nevertheless, the behavior will remain the same. We can see that Page-Consumer received only 2 events, and the rest were filtered out by OnlinePageProducerConsumer. We have successfully added another stage with minimal changes.

There is still room for improvement. We kept this example simple on purpose to show you how to extend your pipeline with additional stages. We used a single OnlinePageProducerConsumer process and limited its demand to only 1 event. In production, this could become a performance bottleneck.

To resolve this, we are going to leverage concurrency again by adding another OnlinePageProducerConsumer process to our supervision tree. This will increase our capacity and we will be able to check more pages at the same time.

Scaling Up a Stage With Extra Processes

Since we are adding more processes of the same type, let's create a Registry and start keeping track of them. We can add it as the first child process in our main supervision tree in application.ex:

```
children = [
  {Registry, keys: :unique, name: ProducerConsumerRegistry},
  PageProducer,
  OnlinePageProducerConsumer,
  PageConsumerSupervisor
]
```

If you want to refresh your memory on how to use Registry, please see Naming Processes Using the Registry, on page 49.

We called the new process ProducerConsumerRegistry to keep its name short, but you can also call it OnlinePageProducerConsumerRegistry if you'd like to be specific.

Next, we want to add another OnlinePageProducerConsumer, but we have to assign it a unique name and an id before adding it to the main supervision tree. We may want to add more of these processes in the future, so let's create a helper function to avoid repetitive code:

```
scraper/lib/scraper/application.change4.ex
def producer_consumer_spec(id: id) do
  id = "online_page_producer_consumer_#{id}"
  Supervisor.child_spec({OnlinePageProducerConsumer, id}, id: id)
end
```

The new producer_consumer_spec/1 helper accepts an id param, which we are going to prefix with the string online_page_producer_consumer_ for readability. Notice that the id is used as the unique process id when calling child_spec/2. It is also given to the process as an argument. You will see how we are going to use it in just a moment.

Now, let's use our helper function and see how it works:

```
scraper/lib/scraper/application.change4.ex
children = [
  {Registry, keys: :unique, name: ProducerConsumerRegistry},
  PageProducer,
➤ producer_consumer_spec(id: 1),
➤ producer_consumer_spec(id: 2),
  PageConsumerSupervisor
]
```

We used plain integers for the id param. Calling producer_consumer_spec(id: 1) will therefore return a process specification with an id online_page_producer_consumer_1. The same id is now available to us in the start_link/1 function for OnlinePageProducerConsumer:

```
scraper/lib/online_page_producer_consumer.change1.ex
def start_link(id) do
  initial_state = []
  GenStage.start_link(__MODULE__, initial_state, name: via(id))
end
```

Here's where the id comes in useful—we want to use the :via method of getting a process name and assign it to the process. This is all done in the via/1 helper function, which looks like so:

```
scraper/lib/online_page_producer_consumer.change1.ex
def via(id) do
  {:via, Registry, {ProducerConsumerRegistry, id}}
end
```

We use the ProducerConsumerRegistry we created earlier to store a reference to our process.

We're almost there. The final piece of the puzzle is updating our PageConsumerSupervisor—it needs to know which producers to subscribe to for events. The subscribe_to option for consumer stages is very flexible. We already know the short syntax that accepts a module name. You can also pass a process identifier, or use a :via tuple. We are going to use the latter, since we already have the via/1 helper defined in OnlinePageProducerConsumer:

```
scraper/lib/page_consumer_supervisor.change2.ex
subscribe_to: [
  OnlinePageProducerConsumer.via("online_page_producer_consumer_1"),
  OnlinePageProducerConsumer.via("online_page_producer_consumer_2"),
]
```

Now we have everything in place and we can give it a try by firing iex with the same test data:

```
iex(1)> PageProducer.scrape_pages(pages)
```

```
13:27:44.813 [info]  OnlinePageProducerConsumer received ["google.com"]
13:27:44.813 [info]  OnlinePageProducerConsumer received ["facebook.com"]
13:27:46.814 [info]  PageConsumer received facebook.com
13:27:47.815 [info]  OnlinePageProducerConsumer received ["apple.com"]
13:27:48.816 [info]  PageConsumer received apple.com
13:27:49.814 [info]  PageConsumer received google.com
13:27:51.815 [info]  OnlinePageProducerConsumer received ["amazon.com"]
13:27:51.817 [info]  OnlinePageProducerConsumer received ["netflix.com"]
13:27:52.816 [info]  PageConsumer received amazon.com
13:27:52.816 [info]  Received demand for 1 pages
13:27:52.818 [info]  Received demand for 1 pages
```

In the output above, we can see that both OnlinePageProducerConsumer processes picked up a website each. This helps us process an extra item while the other process is busy, so we effectively doubled our pipeline capacity for this stage of the pipeline with only a few changes. We can also extend this further by adding more OnlinePageProducerConsumer stages. Here is a figure that illustrates what our final data processing pipeline looks like:
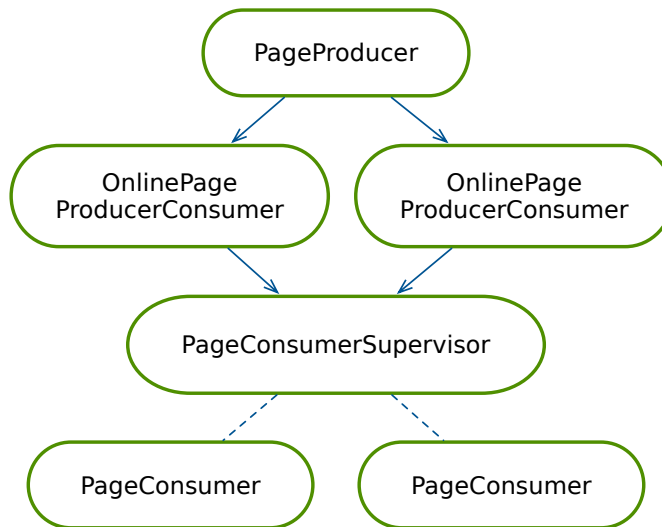
Looking at the figure, you can appreciate the complexity of the system we have just built. Thanks to GenStage, we could focus on the business logic and overall architecture, so we didn't need to write lot of code, which is great.

Hopefully this section helped you understand how :producer_consumer stages help us to build pipelines to tackle any problem. We covered a lot of material, but there is one more thing we're going to discuss, before we wrap up.

Choosing the Right Dispatcher

There is one important component of GenStage that we haven't talked about yet, and this component is the *dispatcher*. When :producer and :producer_consumer stages send events to consumers, it is in fact the dispatcher that takes care of sending the events. So far we have used the default DemandDispatcher, but GenStage comes with two more. Let's see what they do and how we can use them.

You can specify what dispatcher to use when initializing the process, using the :dispatcher key. The default is DemandDispatcher, which is equivalent to this configuration:

```
def init(state) do
  {:producer, state, dispatcher: GenStage.DemandDispatcher}
end
```

DemandDispatcher sends events to consumers with the highest demand first. It is the dispatcher which you're going to use most often. However, there are cases where you may want to route events to consumers using a different strategy, and this is where BroadcastDispatcher and PartitionDispatcher come in. Let's look at BroadcastDispatcher first.

Using BroadcastDispatcher

You can switch to BroadcastDispatcher by setting :dispatcher to GenStage.BroadcastDispatcher in your init/1 function:

```
{:producer, state, dispatcher: GenStage.BroadcastDispatcher}
```

As its name suggests, BroadcastDispatcher sends the events supplied by the :producer or :producer_consumer to *all* consumers subscribed to it. This is useful when you have different types of consumers subscribed to the same producer, and need the same data, but for different purposes.

When BroadcastDispatcher is used, consumer stages get the ability to filter the events they are receiving. This means that each consumer can opt-in for specific events, and discard the rest. All you have to do is use the :selector setting when subscribing to the producer, like so:

```
def init(state) do
  selector =
    fn incoming_event ->
      # you can use the event to decide whether
      # to return `true` and accept it, or `false` to reject it.
    end

  sub_opts = [
```

```
    {SomeProducer, selector: selector}
  ]

  {:consumer, state, subscribe_to: sub_opts}
end
```

As you can see, the :selector is just a function, which you can optionally pass to each producer individually. The consumer is the one that has to examine each message, and decide whether it wants it or not.

The other type of dispatcher we're going to look at also helps you route events to specific consumers, but it works in a slightly different way.

Using PartitionDispatcher

Unlike BroadcastDispatcher, where the consumer has to check each event, PartitionDispatcher leaves this responsibility to the producer. The producer examines each event, and assigns it to a *partition*. You can think of partitions as buckets, which we use to put one or more events in. Consumers can then tell the producer what type events they want to receive, based on that partition.

There are two extra arguments that we need to pass when configuring PartitionDispatcher—:partitions and :hash. Here is an example:

```
def init(state) do
  hash =
    fn event ->
      # you can use the event to decide which partition
      # to assign it to, or use `:none` to ignore it.
      {event, :c}
    end

  opts = [
    partitions: [:a, :b, :c],
    hash: hash
  ]

  {:producer, state, dispatcher: {GenStage.PartitionDispatcher, opts}}
end
```

The hash function is similar to the selector function in BroadcastDispatcher, but it has to return a tuple, rather than a boolean. The tuple should contain the event that we want to dispatch, and the name of the partition. Partitions are chosen by us and could be any Elixir term. Here we picked :a, :b, and :c. You can also ignore an event by passing :none, which you don't have to define in your :partitions list. Also notice that the :dispatcher value is now a tuple, which allows us to pass the dispatcher configuration options as the second element in the tuple.

Now that the producer is configured, consumers can subscribe to one of the partitions when initializing:

```
sub_opts = [
  {SomeProducer, partition: :b}
]
```

```
{:consumer, state, subscribe_to: sub_opts}
```

BroadcastDispatcher and PartitionDispatcher are a great addition to the default DemandDispatcher, and should help you adapt your data processing pipeline to an ever wider variety of uses cases.

However, if they still don't quite match what you are trying to accomplish, you can also create your own dispatcher from scratch, by implementing the GenStage.Dispatcher behaviour. Check out the Dispatcher module documentation[2] for more information on what callbacks you have to implement.

Now it's time for a quick recap.

Wrapping Up

This chapter wasn't the easiest, so congratulations on completing it! Although slightly intimidating at first, GenStage is a brilliant tool that makes other programming languages jealous of Elixir. Of course, as José Valim has pointed out many times, none of it would be possible without Erlang's rock-solid OTP serving as a foundation.

In this chapter we looked at back-pressure and the problems it solves when building and scaling data processing pipelines. You learned about GenStage and how stages enable us to use back-pressure to easily build data processing pipelines. We saw how you can scale your system to handle increased workloads by leveraging concurrency with ConsumerSupervisor. When working on more complex tasks, multi-stage pipelines can help model and tackle potential challenges. You can use and adapt these techniques to build resilient and scalable systems that perform well under increased workloads.

Further reading

You saw that GenStage is very versatile and can be configured in many different ways to solve an endless variety of problems. Although we covered a lot of ground, we have only seen a fraction of what's possible. There are more advanced features in GenStage that didn't make it in this chapter, such as the ability to switch to :manual subscription mode in stages of type :consumer and :producer_consumer. This gives you the flexibility to

---

2. https://hexdocs.pm/gen_stage/GenStage.Dispatcher.html#content

change demand at runtime, rather than hard-coding `max_demand` and `min_demand` at compile time.

These and other features are well documented online, so I would encourage you to check GenStage's official documentation[a] and explore further what's available.

_____

a.   https://hexdocs.pm/gen_stage/GenStage.html

Working with GenStage could be challenging. In order for your data processing pipeline to work, everything has to be configured and set up correctly. To achieve the desired level of concurrency, you may have to create more than a few types of stages. Often this makes GenStage a less appealing choice for more common tasks and problems, especially in teams less experienced with it. The good news is that José and team have worked on two more libraries—Flow and Broadway—which are built on top of GenStage and are designed to solve specific problems with less configuration and set up. In the next chapter, we're going to learn about Flow, which provides a high-level API for working with collections of data, similar to the Enum and Stream modules. We're going to cover Broadway straight after that in the final chapter.

Processing Collections with Flow

Content to be supplied later.

Data-Ingestion Pipelines with Broadway

Content to be supplied later.

Thank you!

How did you enjoy this book? Please let us know. Take a moment and email us at support@pragprog.com with your feedback. Tell us your story and you could win free ebooks. Please use the subject line "Book Feedback."

Ready for your next great Pragmatic Bookshelf book? Come on over to https://pragprog.com and use the coupon code BUYANOTHER2021 to save 30% on your next ebook.

Void where prohibited, restricted, or otherwise unwelcome. Do not use ebooks near water. If rash persists, see a doctor. Doesn't apply to *The Pragmatic Programmer* ebook because it's older than the Pragmatic Bookshelf itself. Side effects may include increased knowledge and skill, increased marketability, and deep satisfaction. Increase dosage regularly.

And thank you for your continued support.

The Pragmatic Bookshelf

The Pragmatic Bookshelf

The Pragmatic Bookshelf features books written by professional developers for professional developers. The titles continue the well-known Pragmatic Programmer style and continue to garner awards and rave reviews. As development gets more and more difficult, the Pragmatic Programmers will be there with more titles and products to help you stay on top of your game.

Visit Us Online

This Book's Home Page
https://pragprog.com/book/sgdpelixir
Source code from this book, errata, and other resources. Come give us feedback, too!

Keep Up to Date
https://pragprog.com
Join our announcement mailing list (low volume) or follow us on twitter @pragprog for new titles, sales, coupons, hot tips, and more.

New and Noteworthy
https://pragprog.com/news
Check out the latest pragmatic developments, new titles and other offerings.

Buy the Book

If you liked this ebook, perhaps you'd like to have a paper copy of the book. Paperbacks are available from your local independent bookstore and wherever fine books are sold.

Contact Us